

Introduction to Parallel Computing

Shared-memory programming
with OpenMP

Introduction

- OpenMP is an API for shared-memory Multiple Instruction Multiple Data (MIMD) programming.
- OpenMP is designed for systems in which each thread/process has access to all available memory.




FIGURE 5.1

A shared-memory system.

- It allows the programmer to state that a block of code should be executed in parallel, and the precise determination of the tasks and which thread should execute them is left to the compiler and the run-time system.
- Some compilers and languages that support OpenMP:
 - GNU: GCC – C/C++/Fortran
 - LLVM: Clang – C/C++
 - Microsoft: MSVC – C/C++
 - Nvidia: C/C++/Fortran
- OpenMP API is directive-based.
 - *Directives* are special preprocessor instructions called *pragma*.
- The *#pragma* directive is a special purpose directive that is used to turn on or off some features.

Program 5.1: A “hello, world” program that uses OpenMP.

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void hello(void);

int main(int argc, char *argv[]) {
    int thread_count = strtol(argv[1], NULL, 10);

    #pragma omp parallel num_threads(thread_count)
    hello();

    return 0;
}

void hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);
}
```

- *strtol* is a function in the *stdlib.h* header that converts strings to long integer. It's used to convert the command line argument *argv[1]* into a string. The first parameter is string value to be converted, the second is a reference to the end of the string, the third is the base of long value.
- *omp parallel num_threads()* tells the compiler that next line will be executed in parallel by a specific number of threads.
- *omp_get_thread_num()*: a function returns the number (rank) of the current thread.
- *omp_get_num_threads()*: a function returns the total number of threads that are assigned to the process.
- To compile and run the previous program:
 - For Clang and Apple M1

```
clang -Xpreprocessor -fopenmp -I/opt/homebrew/opt/libomp/include -
L/opt/homebrew/opt/libomp/lib -lomp -o my_program <filename.c>
```

- For GCC on Windows, Linux, and Mac

```
gcc-14 -o my_program -fopenmp <filename.c>
```

- Or

```
gcc -o my_program -fopenmp <filename.c>
```

- Then, run using

```
./my_program 4
```

- The program output is different from a run to another. This is because each thread is competing for access to *stdout*.


```
Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 3 of 4

Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4
Hello from thread 3 of 4
```


or

```
Hello from thread 3 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4
```

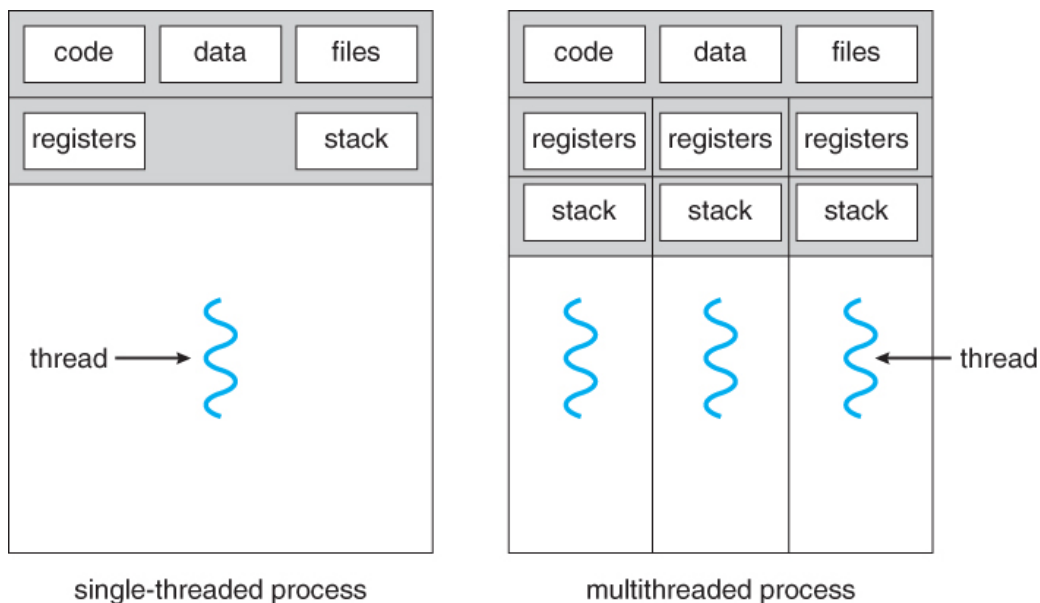
- When the program reaches the *parallel* directive, the original thread forks *thread_count* – 1 additional threads.
- The collection of threads executing the *parallel* block – the original thread and the new threads – is called a **team**.
 - **Master** thread: the first thread of execution, thread 0
 - **Parent** thread: thread that encountered a *parallel* directive and started a team of threads.
 - In many cases, the parent is also the master thread.
 - **Child** thread: each thread started by the parent.



- **Implicit barrier:** a thread that has completed the parallel block of code will wait for all the other threads in the team to complete the block.



- Each thread has its own stack; so, a thread executing the *Hello* function will create its own private, local variables in the function.
- Since *stdout* is shared among the threads, each thread can execute the *printf()* to print its *rand* and the number of threads.
 - There is no scheduling of access to *stdout*, so the actual order in which the threads print their results is **nondeterministic**.



- To handle missing *omp* or if a compiler doesn't support *omp*, we can modify the previous program as follows:

```
#include <stdio.h>
#include <stdlib.h>

#ifdef _OPENMP
#include <omp.h>
#endif

void hello(void);

int main(int argc, char *argv[]) {
    int thread_count = strtol(argv[1], NULL, 10);

    #pragma omp parallel num_threads(thread_count)
    hello();


    return 0;
}

void hello(void) {
#ifdef _OPENMP
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
#else
    int my_rank = 0;
    int thread_count = 1;
#endif

    printf("Hello from thread %d of %d\n", my_rank, thread_count);
}
```

The trapezoidal rule

- The trapezoidal rule is used to estimate the area under the curve.
- To estimate the area between the graph $f(x)$, the vertical lines $x = a$ and $x = b$ and the x-axis, divide the interval $[a, b]$ into n subintervals and approximating the area over each subinterval by the area of a trapezoid.



- Let
 - n is number of subintervals
 - b and a are the limits of the intervals
 - $h = (b - a)/n$
 - $x_i = a + ih$ for $i = 0, 1, \dots, n$

- Then the approximation is

```

/* Input:  a, b, n */
h = (b-a)/n;
approx = (f(a) + f(b))/2.0;
for (i = 1; i <= n-1; i++) {
    x_i = a + i*h;
    approx += f(x_i);
}
approx = h*approx;

```

```

#include <stdio.h>

double f(double x); /* Function we're integrating */
double Trap(double a, double b, int n, double h);

int main(void) {
    double integral; /* Store result in integral */
    double a, b; /* Left and right endpoints */
    int n; /* Number of trapezoids */
    double h; /* Height of trapezoids */

    printf("Enter a, b, and n\n");
    scanf("%lf", &a);
    scanf("%lf", &b);
    scanf("%d", &n);

    h = (b - a) / n;
    integral = Trap(a, b, n, h);

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.15f\n", a, b, integral);

    return 0;
} /* main */

double Trap(double a, double b, int n, double h) {
    double integral = (f(a) + f(b)) / 2.0;
    for (int k = 1; k <= n - 1; k++) {
        integral += f(a + k * h);
    }
    integral = integral * h;

    return integral;
} /* Trap */

double f(double x) {
    return x * x;
} /* f */

```

Parallelizing the trapezoidal rule

Partition the problem solution into tasks.

- Find the area of a single trapezoid
- Computing the sum of these areas

Identify communication channels

- Join the task of computing the area of each trapezoid with the task of computing the final sum

Aggregate tasks into composite tasks

- Assign a contiguous block of trapezoids to each thread.
- Partition the interval $[a, b]$ into larger subintervals.

Map composite tasks to cores

- Each thread applies the serial trapezoidal rule to its subinterval
- Each thread will be executed on a single core





FIGURE 3.5

Tasks and communications for the trapezoidal rule.




```

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>

double f(double x);    /* Function we're integrating */
void trap(double a, double b, int n, double *global_result_p);

int main(int argc, char *argv[]) {
    double global_result = 0.0; /* Store result in global_result */
    double a, b;                /* Left and right endpoints */
    int n;                       /* Total number of trapezoids */

    int thread_count = strtol(argv[1], NULL, 10);
    printf("Enter a, b, and n\n");
    scanf("%lf %lf %d", &a, &b, &n);

    # pragma omp parallel num_threads(thread_count)
        trap(a, b, n, &global_result);

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.14e\n", a, b, global_result);
    return 0;
} /* main */

double f(double x) {
    return x * x;
} /* f */

void trap(double a, double b, int n, double *global_result_p) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
    double x;

    double h = (b - a) / n;
    int local_n = n / thread_count;
    double local_a = a + my_rank * local_n * h;
    double local_b = local_a + local_n * h;
    double my_result = (f(local_a) + f(local_b)) / 2.0;
    for (int i = 1; i <= local_n - 1; i++) {
        x = local_a + i * h;
        my_result += f(x);
    }
    my_result = my_result * h;

    # pragma omp critical
        *global_result_p += my_result;
} /* trap */

```

- The program reads the number of threads to fork as a command line argument.
- Next, the program asks the user to enter: a , the start of the interval, b , the end of the interval, and n the number of subintervals (the smaller trapezoids).
- Notice that when we call the `trap` function, we pass a, b, n as values, while the `global_result` is passed as an address.

- Inside *trap()*:
 - The program gets the rank (ID) of the current thread. This used to compute the start and the end of the interval that is assigned to the current thread.
 - *thread_count* is used to store the total number of threads used to execute the function. This is used to compute how many subintervals (smaller trapezoids) to each thread.
 - *h* is the height value of each trapezoid, which is computed by dividing the size of the interval ($b - a$) by the total number of trapezoids (n).
 - *local_n* is a variable that computes how many trapezoids are assigned to each thread. It is computed by dividing the total number of trapezoids by the total number of threads available.
 - *local_a* and *local_b* store the limits of each subinterval (trapezoid) assigned to each thread.
 - *my_result* is local integral value of each thread. It's computed using the *local_a* and *local_b* instead of the limits of the entire interval $[a, b]$.
 - `#pragma omp critical` is a directive that tells the compiler to make the next statement to have **mutually exclusive** access. This means that the statement `*global_result += my_result` will be executed by only one thread. If more than one thread is executing the same statement at the same time, the result will be incorrect.
 - A code that includes access to a shared resource and causes a **race condition**, is called a **critical section**.
 - **Race condition:** multiple threads are attempting to access a shared resource, at least one of the accesses is an update, and the accesses can result in an error.
- Notice that the variables *local_a*, *local_b*, and *my_result* are different for each thread.
- If number of total trapezoids, n , is not divisible by the number of available threads, *thread_count*, then we will result in inaccurate results and we have to include error handling mechanisms.
- For example, if $n = 14$ and *thread_count* = 4, then each thread will compute

$$local_n = n / thread_count = 14 / 4 = 3$$
 Thus, each thread will use 3 trapezoids and *global_result* will be computed with $4 \times 3 = 12$ instead of 14.
- We can handle this error by including an *if* condition to check that n is divisible by *thread_count*.

```
if ( n % thread_count != 0 ) {
    fprintf (stderr, "n must be evenly divisible by thread_count\n" );
    exit (0) ;
}
```


- During the execution of the program, the $local_a = a + my_rank * local_n * h$; will be as following:

```
thread 0:  a + 0*local_n*h
thread 1:  a + 1*local_n*h
thread 2:  a + 2*local_h*h
. . .
```

- So, the value of my_rank is replaced by the rank of the thread.
- The same applies for the $local_b = local_a + local_n * h$;


Scope of variables

- In serial programming, the **scope** of a variable consists of those parts of a program in which the variable can be used.
- In OpenMP, the **scope** of a variable refers to the set of threads that can access the variable in a parallel block.



Shared variable

all team threads can access it



Private variable

Only one thread can access it

- In the *trapezoidal* program:
 - The variables declared in the **main function** ($a, b, n, global_result, thread_count$) are **shared**; they are declared before the *parallel* directive.
 - The variables inside the **trap** function (h, x, my_rank , etc.) are **private**; they are declared within the scope of parallel directive.
 - The variable **$global_result_p$** in the *trap* function is **private**. However, it refers to the **shared** variable **$global_result$** that is defined in the main function. So, **$global_result_p$** is treated as a **shared variable**.
 - If **$global_result_p$** were **private** to each thread, there would be no need for the *critical* directive. It must be shared to compute the final result computed by individual threads.

- In summary, variables that have been declared before a *parallel* directive have **shared** scope, while variables declared in the block (e.g., local variables in functions) have **private** scope.

The reduction clause

- What is wrong with this code? (it prints the correct result)

```
//gcc-14 -o main -fopenmp p5_5_trapz_parallel_bad.c
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>

double f(double x); /* Function we're integrating */
double local_trap(double a, double b, int n);

int main(int argc, char *argv[]) {
    double global_result = 0.0; /* Store result in global_result */
    double a, b; /* Left and right endpoints */
    int n; /* Total number of trapezoids */

    int thread_count = strtol(argv[1], NULL, 10);
    printf("Enter a, b, and n\n");
    scanf("%lf %lf %d", &a, &b, &n);

    # pragma omp parallel num_threads(thread_count)
    {
        # pragma omp critical
        global_result += local_trap(a, b, n);
    }

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.14e\n", a, b, global_result);
    return 0;
} /* main */

double f(double x) {
    return x * x;
} /* f */

double local_trap(double a, double b, int n) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
    double x;

    double h = (b - a) / n;
    int local_n = n / thread_count;
    double local_a = a + my_rank * local_n * h;
    double local_b = local_a + local_n * h;
    double my_result = (f(local_a) + f(local_b)) / 2.0;
    for (int i = 1; i <= local_n - 1; i++) {
        x = local_a + i * h;
        my_result += f(x);
    }
    my_result = my_result * h;
    return my_result;
} /* local_trap */
```

- The call to *local_trap* can only be executed by one thread at a time.
- Effectively, we're forcing the threads to execute the trapezoidal rule sequentially.
- How to fix the previous code?
 - Just move the critical section after the function call.

```
//gcc-14 -o main -fopenmp p5_6_trapz_parallel_fix.c
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>

double f(double x); /* Function we're integrating */
double local_trap(double a, double b, int n);

int main(int argc, char *argv[]) {
    double global_result = 0.0; /* Store result in global_result */
    double a, b; /* Left and right endpoints */
    int n; /* Total number of trapezoids */
    int thread_count = strtol(argv[1], NULL, 10);
    printf("Enter a, b, and n\n");
    scanf("%lf %lf %d", &a, &b, &n);

    #pragma omp parallel num_threads(thread_count)
    {
        double my_result = 0.0; /* private */
        my_result += local_trap(a, b, n);
    }
    #pragma omp critical
    global_result += my_result;

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.14e\n", a, b, global_result);
    return 0;
} /* main */

double f(double x) {
    return x * x;
} /* f */

double local_trap(double a, double b, int n) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
    double x;
    double h = (b - a) / n;
    int local_n = n / thread_count;
    double local_a = a + my_rank * local_n * h;
    double local_b = local_a + local_n * h;
    double my_result = (f(local_a) + f(local_b)) / 2.0;
    for (int i = 1; i <= local_n - 1; i++) {
        x = local_a + i * h;
        my_result += f(x);
    }
    my_result = my_result * h;
    return my_result;
} /* local_trap */
```

- Now the call to *local_trap* is outside the critical section, and the threads can execute their calls simultaneously.
 - Note that *my_result* is **private** because it is declared in the parallel block.
- OpenMP provides a cleaner method to avoid serializing the execution of *local_trap* by using the **reduction clause**.
- A **reduction clause** consists of two elements:

Reduction operator

- An associative binary operation (+, *)

Reduction variable

- The variable in which the results of an operation (+, *) are stored

Reduction

- A computation that repeatedly applies the same **reduction operator** to a **sequence of operands** to get a single result

- For example, if *A* is an array of *n* integers, the computation

```
int sum = 0;
for (i = 0; i < n; i++)
    sum += A[i];
```

is a reduction in which the **reduction operator** is + and the **reduction variable** is *sum*.

- Reduction clause in OpenMP is defined as follows

```
global_result = 0.0;
# pragma omp parallel num_threads(thread_count) \
    reduction(+: global_result)
global_result += Local_trap(double a, double b, int n);
```

- < *parallel directive* > *reduction*(< *operator* >: < *variable list* >)
- Reduction variable → *global_result*
- Reduction operator → +

```

//gcc-14 -o main -fopenmp p5_7_trapz_parallel_reduction.c
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

double f(double x); /* Function we're integrating */
double local_trap(double a, double b, int n);

int main(int argc, char *argv[]) {
    double global_result = 0.0; /* Store result in global_result */
    double a, b; /* Left and right endpoints */
    int n; /* Total number of trapezoids */
    int thread_count;

    thread_count = strtol(argv[1], NULL, 10);
    printf("Enter a, b, and n\n");
    scanf("%lf %lf %d", &a, &b, &n);

    # pragma omp parallel num_threads(thread_count) \
        reduction(+: global_result)
        global_result += local_trap(a, b, n);

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.14e\n",
        a, b, global_result);
    return 0;
} /* main */

double f(double x) {
    double return_val;

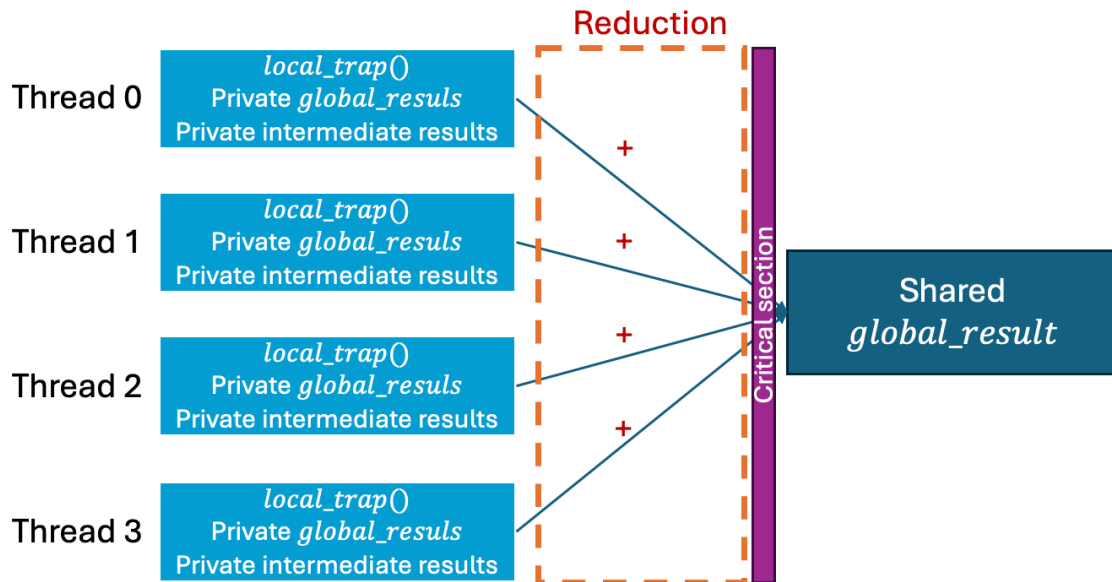
    return_val = x * x;
    return return_val;
} /* f */

double local_trap(double a, double b, int n) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
    double x;
    double h = (b - a) / n;
    int local_n = n / thread_count;
    double local_a = a + my_rank * local_n * h;
    double local_b = local_a + local_n * h;
    double my_result = (f(local_a) + f(local_b)) / 2.0;
    for (int i = 1; i <= local_n - 1; i++) {
        x = local_a + i * h;
        my_result += f(x);
    }
    my_result = my_result * h;

    return my_result;
} /* Trap */

```

- Behind the scenes,
 - OpenMP creates a private variable for each thread
 - The run-time system stores each thread's result in this private variable
 - OpenMP creates a critical section, and the values stored in the private variables are added in this critical section.



- A reduction operator can be `+`, `-`, `*`, `&`, `|`, `^`, `&&` and `||`.
 - Why is **division** not supported? Because it is **not commutative** or **associative**.
 - Although **subtraction** is not commutative or associative, it's handled in a different way in OpenMP.
- *float* and *double* reduction variables may cause the results to be slightly different.
 - Floating point arithmetic isn't associative.
 - For example, if a , b , and c are floats, then $(a + b) + c \neq a + (b + c)$.
- The private variables of each thread are initialized to a default value according to the datatype of the reduction variable

Table 5.1 Identity values for the various reduction operators in OpenMP.

Operator	Identity Value
<code>+</code>	0
<code>*</code>	1
<code>-</code>	0
<code>&</code>	<code>~0</code>
<code> </code>	0
<code>^</code>	0
<code>&&</code>	1
<code> </code>	0

The *parallel for* directive

- Instead of using explicit parallelization, we can use the *parallel for* directive
- It's placed immediately before the *for* loop.
 - Serial program vs parallel program

```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```

```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
# pragma omp parallel for num_threads(thread_count) \  
    reduction(+: approx)  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```

- The system divides the iterations of the *for* loop among the threads.
- The loop variable *i* has a default **private** scope; so each thread has its own copy of *i*.

```
#include <stdio.h>  
  
double f(double x);  
double Trap(double a, double b, int n, double h);  
  
int main(void) {  
    double integral;  
    double a, b;  
    int n;  
    double h;  
  
    printf("Enter a, b, and n\n");  
    scanf("%lf", &a);  
    scanf("%lf", &b);  
    scanf("%d", &n);  
  
    h = (b - a) / n;  
    integral = Trap(a, b, n, h);  
  
    printf("With n = %d trapezoids, our estimate\n", n);  
    printf("of the integral from %f to %f = %.15f\n", a, b, integral);  
  
    return 0;  
} /* main */  
  
double Trap(double a, double b, int n, double h) {  
    double integral = (f(a) + f(b)) / 2.0;  
  
    #pragma omp parallel for num_threads(4) reduction(+:integral)  
    for (int k = 1; k <= n - 1; k++) {  
        integral += f(a + k * h);  
    }  
    integral = integral * h;  
  
    return integral;  
} /* Trap */  
  
double f(double x) {  
    return x * x;  
} /* f */
```

- Note the following:
 1. OpenMP only parallelize *for* loop – no parallelization for *while* and *do – while* loops
 2. The number of iterations of the *for* loop must be determined
 - a. Infinite loops cannot be parallelized

```
for ( ; ; ) {
    . . .
}
```

- b. *for* loops with *break* cannot be parallelized

```
for (i = 0; i < n; i++) {
    if ( . . . ) break;
    . . .
}
```

- Notice the error in this code

```
#include <stdio.h>
#include <omp.h>

int search(int key, int arr[], int n) {
    int i;
    #pragma omp parallel for
    for (i = 0; i < 5; i++) {
        if (key == arr[i])
            return i;
    }
    return -1;
}

int main(int argc, char *argv[]) {
    int n;
    int arr[] = {1, 2, 3, 4, 5};
    int key = 4;

    scanf("%d", &n);
    int res = search(key, arr, n);
    printf("%d\n", res);
    return 0;
}
```

- The error is at the *return i;* statement.
- OpenMP cannot **exit** a loop that is parallelized.

- Try this Fibonacci program without and with parallelization:

```
#include <stdio.h>
#include <omp.h>

int main(int argc, char *argv[]) {
    int fib[20];
    fib[0] = fib[1] = 1;

    #pragma omp parallel for
    for (int i = 2; i < 20; ++i) {
        fib[i] = fib[i - 1] + fib[i - 2];
    }

    for (int i = 0; i < 20; ++i) {
        printf("%d ", fib[i]);
    }
    return 0;
}
```

- When parallelization is used, the output is unpredictable.
- This is because the Fibonacci program includes **dependencies** – each iteration is dependent on the previous iterations. Hence, **we cannot compute the value at a specific index without computing the value at lower indices.**
- In summary, OpenMP don't check for dependences among iterations in a loop.

Estimating π

- To compute a numerical approximation to π , we use the following formula

$$\pi = 4 \left[1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots \right] = 4 \sum_{k=0}^{\infty} \frac{(-1)^k}{2k+1}.$$

- The serial program

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    double factor = 1.0;
    double sum = 0.0;
    for (int i = 0; i < 1000; ++i) {
        sum += factor / (2 * i + 1);
        factor = -factor;
    }
    double pi = 4 * sum;
    printf("PI = %lf", pi);
    return 0;
}
```

- The parallel program

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    double factor = 1.0;
    double sum = 0.0;
    #pragma omp parallel for
    for (int i = 0; i < 1000; ++i) {
        sum += factor / (2 * i + 1);
        factor = -factor;
    }
    double pi = 4 * sum;
    printf("PI = %lf", pi);
    return 0;
}
```

- This program has a problem: there is a **data dependency** between *factor* and *sum*, thus it gives **incorrect** results.
- We can fix the previous program by replacing

```
sum += factor / (2 * k + 1);
factor = -factor;
```

by

```
if (k % 2 == 0)
    factor = 1.0;
else
    factor = -1.0;
sum += factor / (2 * k + 1);
```

or

```
factor = (k % 2 == 0) ? 1.0 : -1.0;
sum += factor / (2 * k + 1);
```

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    double factor = 1.0;
    double sum = 0.0;
    #pragma omp parallel for
    for (int i = 0; i < 1000; ++i) {
        factor = (i % 2 == 0) ? 1.0 : -1.0;
        sum += factor / (2 * i + 1);
    }
    double pi = 4 * sum;
    printf("PI = %lf", pi);
    return 0;
}
```

- This program gives correct results, but not accurate!

- To fix the previous program:
 1. We must tell OpenMP that *sum* is a **reduction** variable, thus the final result is accumulated into it
 2. The *factor* variable is shared because it's defined before the *parallel for* directive. We must tell OpenMP that *factor* is private.

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    double factor = 1.0;
    double sum = 0.0;
    #pragma omp parallel for reduction(+:sum) private(factor)
    for (int i = 0; i < 1000; ++i) {
        factor = (i % 2 == 0) ? 1.0 : -1.0;
        sum += factor / (2 * i + 1);
    }
    double pi = 4 * sum;
    printf("PI = %lf", pi);
    return 0;
}
```

- Instead of letting OpenMP to choose the scope of the variables, we can explicitly specify the scope using the **default** clause

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    double factor = 1.0;
    double sum = 0.0;
    int i;
    int n = 1000;
    #pragma omp parallel for default(none) reduction(+:sum)\
    private(i, factor) shared(n)
    for (i = 0; i < n; ++i) {
        factor = (i % 2 == 0) ? 1.0 : -1.0;
        sum += factor / (2 * i + 1);
    }
    double pi = 4 * sum;
    printf("PI = %lf", pi);
    return 0;
}
```

Sorting

Bubble sort

- The serial bubble sort algorithm:

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    int n = 10;
    int a[] = {10, -5, 9, 1, 0, 2, 4, 3, 6, 11};

    for (int i = n-1; i > 0; i--) {
        for (int j = 0; j < i; j++) {
            if (a[j] > a[j + 1]) {
                int tmp = a[j];
                a[j] = a[j + 1];
                a[j + 1] = tmp;
            }
        }
    }

    for (int i = 0; i < n; i++) {
        printf("a[%d] = %d\n", i, a[i]);
    }
    return 0;
}
```

- The outer loop handles all the n-element array
- The inner loop handles the first n-1 element array
- The inner loop compares consecutive pairs of elements in the current list. When a pair is out of order ($a[i] > a[i + 1]$) it swaps them.
- **The inner loop depends on the outer loop.**
- **The inner loop itself also has dependence on previous iterations:**
 - Say an iteration j will swap the values at index $a[j]$ and $a[j + 1]$, this will impact iteration $j + 1$.
 - So, iteration $j + 1$ cannot decide whether to swap the elements or not, until iteration j finishes.
- We cannot remove the **loop-carried dependence** without completely rewriting the algorithm.

Odd-even transposition sort

- A sorting algorithm similar to bubble sort but can be parallelized.
- The serial algorithm is as follows

```
for (phase = 0; phase < n; phase++)
    if (phase % 2 == 0)
        for (i = 1; i < n; i += 2)
            if (a[i-1] > a[i]) Swap(&a[i-1], &a[i]);
    else
        for (i = 1; i < n-1; i += 2)
            if (a[i] > a[i+1]) Swap(&a[i], &a[i+1]);
```

Table 5.2 Serial odd-even transposition sort.

Phase	Subscript in Array				
	0		1	2	3
0	9	↔	7	8	↔ 6
	7		9	6	8
1	7		9	↔ 6	8
	7		6	9	8
2	7	↔ 6		9	↔ 8
	6		7	8	9
3	6		7	↔ 8	9
	6		7	8	9

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    int phase, i, tmp;
    int n = 10;
    int a[] = {10, -5, 9, 1, 0, 2, 4, 3, 6, 11};

    for (phase = 0; phase < n; phase++) {
        if (phase % 2 == 0) {
            for (i = 1; i < n; i += 2) {
                if (a[i - 1] > a[i]) {
                    tmp = a[i - 1];
                    a[i - 1] = a[i];
                    a[i] = tmp;
                }
            }
        } else {
            for (i = 1; i < n - 1; i += 2) {
                if (a[i] > a[i + 1]) {
                    tmp = a[i + 1];
                    a[i + 1] = a[i];
                    a[i] = tmp;
                }
            }
        }
        for (i = 0; i < n; i++) {
            printf("a[%d] = %d\n", i, a[i]);
        }
    }
    return 0;
}
```

- There is a **carried dependence** in the outer loop; two iterations cannot be executed simultaneously.
- The inner loops do not have dependence; each of the inner loops can be executed simultaneously.
 - This because, say the even-phase for loop, compares two adjacent elements. So, for two distinct values of i , say $i = j$ and $i = k$, the pairs $\{j - 1, j\}$ and $\{k - 1, k\}$ will be disjoint.
 - Hence, the comparison and possible swaps of the pairs $(a[j - 1], a[j])$ and $(a[k - 1], a[k])$ can proceed simultaneously.

```

#include <stdio.h>

int main(int argc, char *argv[]) {
    int phase, i, tmp;
    int n = 10;
    int a[] = {10, -5, 9, 1, 0, 2, 4, 3, 6, 11};

    for (phase = 0; phase < n; phase++) {
        if (phase % 2 == 0) {
            #pragma omp parallel for num_threads(4) shared(a, n) private(i, tmp)
            for (i = 1; i < n; i += 2) {
                if (a[i - 1] > a[i]) {
                    tmp = a[i - 1];
                    a[i - 1] = a[i];
                    a[i] = tmp;
                }
            }
        } else {
            #pragma omp parallel for num_threads(4) shared(a, n) private(i, tmp)
            for (i = 1; i < n - 1; i += 2) {
                if (a[i] > a[i + 1]) {
                    tmp = a[i + 1];
                    a[i + 1] = a[i];
                    a[i] = tmp;
                }
            }
        }
        for (i = 0; i < n; i++) {
            printf("a[%d] = %d\n", i, a[i]);
        }
        return 0;
    }
}

```

- Compare the performance of the bubble sort vs the odd-even sort on an array of 100,000 integers.
 - *Comment either the bubble sort or the odd-even sort block, when testing.*
- You will notice that the parallel sort algorithm is almost 2x faster than the serial one.


```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define ARRAY_SIZE 100000

void generate_random_array(int *array, int size) {
    // seed the random number generator
    srand(time(NULL));

    for (int i = 0; i < size; i++) {
        array[i] = rand();
    }
}

void bubble_sort(int *array, int size) {
    for (int i = size - 1; i > 0; i--) {
        for (int j = 0; j < i; j++) {
            if (array[j] > array[j + 1]) {
                // Compare adjacent elements
                int tmp = array[j];
                array[j] = array[j + 1];
                array[j + 1] = tmp;
            }
        }
    }
}

void odd_even_sort(int *array, int size) {
    int phase, i, tmp;

    for (phase = 0; phase < size; phase++) {
        if (phase % 2 == 0) {
            #pragma omp parallel for num_threads(4) shared(array, size) private(i, tmp)
            for (i = 1; i < size; i += 2) {
                if (array[i - 1] > array[i]) {
                    tmp = array[i - 1];
                    array[i - 1] = array[i];
                    array[i] = tmp;
                }
            }
        } else {
            #pragma omp parallel for num_threads(4) shared(array, size) private(i, tmp)
            for (i = 1; i < size - 1; i += 2) {
                if (array[i] > array[i + 1]) {
                    tmp = array[i + 1];
                    array[i + 1] = array[i];
                    array[i] = tmp;
                }
            }
        }
    }
}

int main(int argc, char *argv[]) {
    int array[ARRAY_SIZE];
    // Generate the random array
    generate_random_array(array, ARRAY_SIZE);

    time_t start, end;

    time(&start);
    bubble_sort(array, ARRAY_SIZE);
    time(&end);
    printf("%f\n", difftime(end, start));

    time(&start);
    odd_even_sort(array, ARRAY_SIZE);
    time(&end);
    printf("%f\n", difftime(end, start));

    return 0;
}

```

- The previous parallel odd-even sort has one issue: omp forks the threads at the beginning of the first *for* loop and joins them at the end. The team of threads is forked at the beginning of the second *for* loop and joined again at the end.
- This fork-join process costs time.
- We can solve this issue by forking the threads once at the beginning of the outer loop and only join them at the end of the second inner loop.

```
#define ARRAY_SIZE 200000

void enhanced_odd_even_sort(int *array, int size) {
    int phase, i, tmp;
    #pragma omp parallel num_threads(4) shared(array, size)
    private(i, tmp, phase)
        for (phase = 0; phase < size; phase++) {
            if (phase % 2 == 0) {
                #pragma omp for
                for (i = 1; i < size; i += 2) {
                    if (array[i - 1] > array[i]) {
                        tmp = array[i - 1];
                        array[i - 1] = array[i];
                        array[i] = tmp;
                    }
                }
            } else {
                #pragma omp for
                for (i = 1; i < size - 1; i += 2) {
                    if (array[i] > array[i + 1]) {
                        tmp = array[i + 1];
                        array[i + 1] = array[i];
                        array[i] = tmp;
                    }
                }
            }
        }
}
```

- Run this enhanced odd-even sort in comparison with the basic odd-sort function on an array of size 200,000.
- You will notice that this new function is slightly faster than the basic function.
- In the enhanced function:
 1. We defined the outer loop as a parallel structure, but without using *parallel for* directive. So, this will fork the threads.
 2. At the beginning of each of the inner loops, we used the *for* directive. This will let the loops to execute in parallel using the forked threads, without forking another team of threads.

Scheduling loops

- The *parallel for* directive automatically applies **block partitioning**: if the serial loop has n iterations, then thread 0 handles the first $n/\text{thread_count}$ iterations, thread 1 handles the second $n/\text{thread_count}$ iterations, and so on.
- This could be **less optimal**.
- For example, say we want to parallelize the serial loop

```
sum = 0.0;
for (i = 0; i <= n; i++)
    sum += f(i);
```

Where the time required to compute f is **proportional** to the size of the argument i .

- Then, block partitioning will assign much more work to the last thread, $\text{thread_count} - 1$, than it will assign to thread 0
- It's better to use **cyclic partitioning** for better **work-balance** between threads: the iterations are assigned, one at a time, in a “round-robin” fashion to the threads.
- Suppose $t = \text{thread_count}$. Then a **cyclic partitioning** will assign the iterations as follows

Thread	Iterations
0	0, n/t , $2n/t$, ...
1	1, $n/t + 1$, $2n/t + 1$, ...
\vdots	\vdots
$t - 1$	$t - 1$, $n/t + t - 1$, $2n/t + t - 1$, ...

- Example:

```
double f(int i) {
    int j, start = i*(i+1)/2, finish = start + i;
    double return_val = 0.0;

    for (j = start; j <= finish; j++) {
        return_val += sin(j);
    }
    return return_val;
} /* f */
```

- Run the next three programs and compare the performance.
`gcc-14 -o prog -fopenmp <filename.c>; ./prog`

- The serial program

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>

double Sum(long n);
double f(long i);

int main(int argc, char* argv[]) {
    double global_result;
    double start, finish;
    long n = 100000;

    start = omp_get_wtime();
    global_result = Sum(n);
    finish = omp_get_wtime();

    printf("Result = %.14e\n", global_result);
    printf("Elapsed time = %f seconds\n", finish-start);

    return 0;
}

double f(long i) {
    long j;
    long start = i*(i+1)/2;
    long finish = start + i;
    double return_val = 0.0;

    for (j = start; j <= finish; j++) {
        return_val += sin(j);
    }
    return return_val;
}

double Sum(long n) {
    double approx = 0.0;
    long i;

    for (i = 0; i <= n; i++) {
        approx += f(i);
    }
    return approx;
}
```

- The parallel program

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>

double Sum(long n, int thread_count);
double f(long i);

int main(int argc, char* argv[]) {
    double global_result;
    double start, finish;

    int thread_count = 2;
    long n = 100000;

    start = omp_get_wtime();
    global_result = Sum(n, thread_count);
    finish = omp_get_wtime();

    printf("Result = %.14e\n", global_result);
    printf("Elapsed time = %f seconds\n", finish-start);

    return 0;
}

double f(long i) {
    long j;
    long start = i*(i+1)/2;
    long finish = start + i;
    double return_val = 0.0;

    for (j = start; j <= finish; j++) {
        return_val += sin(j);
    }
    return return_val;
}

double Sum(long n, int thread_count) {
    double approx = 0.0;
    long i;

    # pragma omp parallel for num_threads(thread_count) \
    reduction(+: approx)
    for (i = 0; i <= n; i++) {
        approx += f(i);
    }
    return approx;
}
```

- The **scheduled** parallel program

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>

double Sum(long n, int thread_count);
double f(long i);

int main(int argc, char *argv[]) {
    double global_result;
    double start, finish;

    int thread_count = 2;
    long n = 100000;

    start = omp_get_wtime();
    global_result = Sum(n, thread_count);
    finish = omp_get_wtime();

    printf("Result = %.14e\n", global_result);
    printf("Elapsed time = %f seconds\n", finish - start);

    return 0;
}

double f(long i) {
    long j;
    long start = i * (i + 1) / 2;
    long finish = start + i;
    double return_val = 0.0;

    for (j = start; j <= finish; j++) {
        return_val += sin(j);
    }
    return return_val;
}

double Sum(long n, int thread_count) {
    double approx = 0.0;
    long i;

    #pragma omp parallel for num_threads(thread_count) \
    reduction(+: approx) schedule(dynamic)
    for (i = 0; i <= n; i++) {
        approx += f(i);
    }
    return approx;
}
```

- **Scheduling** in OpenMP: assigning iterations to threads.
- The ***schedule* clause** can be used to assign iterations to threads:

$$\text{schedule}(< \text{type} > [, \text{chunk size}])$$
- **Chunk**: a block of iterations that would be executed consecutively in the serial loop.
 - *chunk size* determines the **number of iterations in a block**. It's **optional**.

- Types of scheduling:

static

- The iterations are assigned to the threads before the loop is executed

dynamic/guided

- The iterations are assigned to the threads while the loop is executing
- After a thread completes its current set of iterations, it requests more from the run-time system

auto

- The compiler and/or the run-time system determine the schedule
- *Does not specify chunk size*

runtime

- The schedule is determined at run-time based on an environment variable

The *static* schedule type

- The system assigns **chunks** of ***chunksize* iterations** to each thread in a round-robin fashion.
- Example: suppose we have 12 iterations, 0, 1,...,11, and three threads.
 - For *schedule(static, 1)*, the iterations will be assigned as

```
Thread 0 : 0, 3, 6, 9
Thread 1 : 1, 4, 7, 10
Thread 2 : 2, 5, 8, 11
```

- For *schedule(static, 2)*, the iterations will be assigned as

```
Thread 0 : 0, 1, 6, 7
Thread 1 : 2, 3, 8, 9
Thread 2 : 4, 5, 10, 11
```

- For *schedule(static, 4)*, the iterations will be assigned as

```
Thread 0 : 0, 1, 2, 3
Thread 1 : 4, 5, 6, 7
Thread 2 : 8, 9, 10, 11
```

- The **default** schedule is defined by a particular implementation of OpenMP, but in most cases it is equivalent to the clause:

*schedule(**static** , total_iterations / thread_count)*

- The *static* schedule is a good choice when each loop iteration takes roughly the **same amount of time** to compute.

The *dynamic* and *guided* schedule type


- *dynamic* schedule: the iterations are also **broken up into chunks** of *chunksize* consecutive iterations.
 - Each thread executes a chunk, and when a thread finishes a chunk, it requests another one from the run-time system.
- The **default *chunksize*** is 1.
- The **difference** between *static* and *dynamic* schedules:
 - *dynamic* schedule assigns ranges to threads on a **first-come, first-served** basis.
 - Good choice if loop iterations **do not take a uniform** amount of time to compute.
 - There is some **overhead** associated with assigning them dynamically at run-time.
 - *static* schedule assigns a range of threads **fixedly** at the start of the loop
 - Used when loop iterations **take a uniform** amount of time.
- With **larger** chunk sizes, **fewer** dynamic assignments will be made.
- In a ***guided* schedule**, as chunks are completed, the size of the new chunks **decreases**.
- If no *chunksize* is specified, the size of the chunks decreases down to **1**. If *chunksize* is specified, it decreases down to ***chunksize***.

The *runtime* schedule type


- When *schedule(runtime)* is specified, the system uses the **environment variable OMP_SCHEDULE** to determine at run-time how to schedule the loop.
 - Environment variables are **named values** that can be accessed by a running program.
 - Examples: *PATH, HOME, SHELL*
- The **OMP_SCHEDULE** environment variable can take on any of the values that can be used for a static, dynamic, or guided schedule.
 - `$ export OMP_SCHEDULE="static ,1"`

- Scheduling **visualization** for the **static**, **dynamic**, and **guided** schedule types with **4 threads and 32 iterations**.


`schedule(static)`




`schedule(static, 2)`



`schedule(dynamic, 2)`



`schedule(guided)`




Which schedule type to choose?

- It depends on the problem and the nature of the loop.
- There is some **overhead** associated with the use of a **schedule clause**.
 - *static* overhead **less than** *dynamic* overhead **less than** *guided* overhead
- Practically, you should **try different options** until you hit a satisfactory performance.
 - Make use of the *schedule(runtime)* to test different settings.

Producers and consumers


- A **queue** is a data structure in which elements are inserted at the rear of the queue and removed from the front of the queue.
 - Adding elements → enqueue
 - Removing elements → dequeue



- Queues appear in many **applications**:
 - Print spooling
 - Operating systems scheduling/task scheduling
 - Data buffers in networking
 - Load balancing in web servers
- Queues are fundamental in **producer-consumer** applications:
 - Part the of the application is **producing** data, and another part is **consuming** the data

Producer Threads

Consumer Threads



- Let's implement a **message-passing** application that is based on the **producer-consumer** pattern.
 - Each thread has a **shared-message queue**.
 - Each thread generates random (integer values) messages and random destination for the messages.
 - A **producer** thread **enqueues** the message in the queue.
 - A **consumer** thread checks if the queue received a message, and then **dequeues** it.
 - Each thread **alternates** between sending and receiving messages.
 - The user will specify the **number of messages** each thread should send.
 - When a thread is **done sending messages**, it receives messages until all the threads are done, at which point all the threads **quit**.
- Pseudocode for each thread might look something like this:

```

for ( sent_msgs = 0; sent_msgs < send_max; sent_msgs++) {
    Send_msg ();
    Try_receive ();
}

while (! Done ())
    Try_receive ();

```

- Core segments/functions of the program:
 - Sending messages → *Send_msg()*
 - Receiving messages → *Try_receive()*
 - Terminating the program → *Done()*

Sending messages

- Accessing a **message queue** to **enqueue** a message is a **critical section**.
 - We need to regulate how threads access the queues to avoid dropping messages.
- To successfully **enqueue** a message, we need a pointer to the **rear** of the queue.
 - When we enqueue a message, we'll update the rear pointer.

```

    msg = random ();
    dest = random () % thread_count;
# pragma omp critical
    Enqueue(queue, dest, my_rank, msg);

```

- The *Enqueue* functions handles updating the rear pointer.
- Note that this allows a thread to send a message to itself.

Receiving messages

- **Only the owner of the queue (that is, the destination thread) will dequeue from a given message queue.**
- If there are at least two messages in the queue, a call to *Dequeue* **can't conflict** with any calls to *Enqueue*.
- Before *Dequeueing* the queue, **we must check its size**; if 0, return none. If the size is ≥ 1 , then return the element at the front.
 - We store two variables: *enqueued*, which is the number of added elements, and *dequeued*, which is the number of removed elements.

```
queue_size = enqueued - dequeued;  
if (queue_size == 0) return;  
else if (queue_size == 1)  
#    pragma omp critical  
    Dequeue(queue, &src, &mesg);  
else  
    Dequeue(queue, &src, &mesg);  
    Print_message(src, mesg);
```

- The *critical section* is added to prevent other threads from updating the queue when the owner thread is *Dequeueing* it when it has only one element.

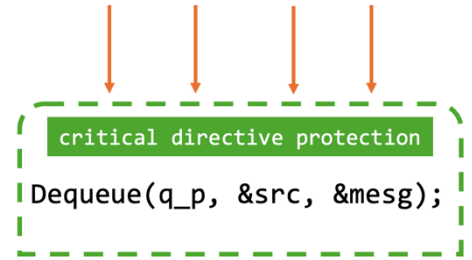
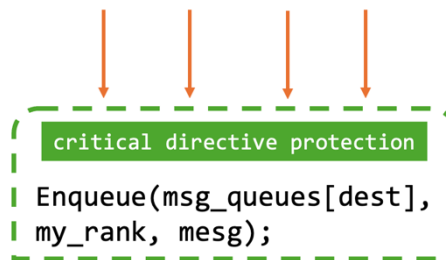
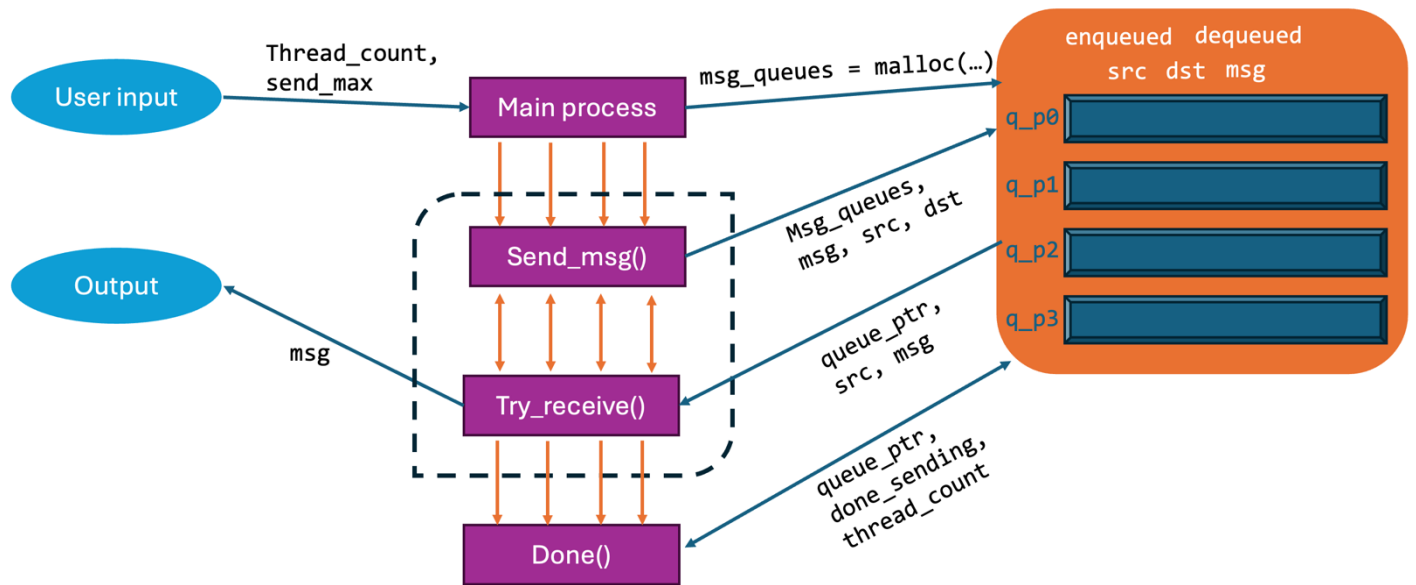
Terminating the program

- The **naïve way** to terminate the program (implementing *Done()*) is as follows:

```
queue_size = enqueued - dequeued;  
if (queue_size == 0)  
    return TRUE;  
else  
    return FALSE;
```

- The problem with this way is that it is possible that the owner thread compute $queue_size = 0$, while another thread is adding an element into the queue.
 - Thus, the added message will never be received by the owner thread.
- Instead, we can include *done_sending* as a counter variable, so it serves as a flag when all threads have no more messages to send.

```
queue_size = enqueued - dequeued;  
if (queue_size == 0 && done_sending == thread_count)  
    return TRUE;  
else  
    return FALSE;
```



```

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include "queue.h"
// gcc-14 -o prog -fopenmp -DDEBUG omp_msgps.c queue.c
const int MAX_MSG = 10000;

void Usage(char* prog_name);
void Send_msg(struct queue_s* msg_queues[], int my_rank, int
thread_count, int msg_number);
void Try_receive(struct queue_s* q_p, int my_rank);
int Done(struct queue_s* q_p, int done_sending, int thread_count);

int main(int argc, char* argv[]) {
    int thread_count;
    int send_max;
    struct queue_s** msg_queues;
    int done_sending = 0;

    if (argc != 3) Usage(argv[0]);
    thread_count = strtol(argv[1], NULL, 10);
    send_max = strtol(argv[2], NULL, 10);
    if (thread_count <= 0 || send_max < 0) Usage(argv[0]);

    msg_queues = malloc(thread_count*sizeof(struct queue_node_s*));

    # pragma omp parallel num_threads(thread_count) \
        default(none) shared(thread_count, send_max, msg_queues,
done_sending)
    {
        int my_rank = omp_get_thread_num();
        int msg_number;
        srandom(my_rank);
        msg_queues[my_rank] = Allocate_queue();

    # pragma omp barrier

        for (msg_number = 0; msg_number < send_max; msg_number++) {
            Send_msg(msg_queues, my_rank, thread_count, msg_number);
            Try_receive(msg_queues[my_rank], my_rank);
        }

    # pragma omp atomic
        done_sending++;
    # ifdef DEBUG
        printf("Thread %d > done sending\n", my_rank);
    # endif

        while (!Done(msg_queues[my_rank], done_sending, thread_count))
            Try_receive(msg_queues[my_rank], my_rank);

        Free_queue(msg_queues[my_rank]);
        free(msg_queues[my_rank]);
    } /* omp parallel */

    free(msg_queues);
    return 0;
} /* main */

```

```

void Usage(char *prog_name) {
    fprintf(stderr, "usage: %s <number of threads> <number of
messages>\n", prog_name);
    fprintf(stderr, "    number of messages = number sent by each
thread\n");
    exit(0);
} /* Usage */

void Send_msg(struct queue_s* msg_queues[], int my_rank, int
thread_count, int msg_number) {
    int mesg = -msg_number;
    int dest = random() % thread_count;
    # pragma omp critical
    Enqueue(msg_queues[dest], my_rank, mesg);
    # ifdef DEBUG
    printf("Thread %d > sent %d to %d\n", my_rank, mesg, dest);
    # endif
} /* Send_msg */

void Try_receive(struct queue_s* q_p, int my_rank) {
    int src, mesg;
    int queue_size = q_p->enqueued - q_p->dequeued;

    if (queue_size == 0) return;
    else if (queue_size == 1)
    # pragma omp critical
        Dequeue(q_p, &src, &mesg);
    else
        Dequeue(q_p, &src, &mesg);
    printf("Thread %d > received %d from %d\n", my_rank, mesg, src);
} /* Try_receive */

int Done(struct queue_s* q_p, int done_sending, int thread_count) {
    int queue_size = q_p->enqueued - q_p->dequeued;
    if (queue_size == 0 && done_sending == thread_count)
        return 1;
    else
        return 0;
} /* Done */

```

- When the program begins execution, the master thread will get command-line arguments and allocate an array of message queues, one for each thread.
- This array needs to be shared among the threads, since any thread can send to any other thread, and hence any thread can enqueue a message in any of the queues.
- We use *omp barrier* to add an explicit barrier so that no thread starts enqueueing or dequeuing until all threads have constructed their queues to avoid errors.
- After completing its sends, each thread increments *done_sending* before proceeding to its final loop of receives.
 - Incrementing *done_sending* is a **critical section**, and we could protect it with a **critical directive**.
 - Instead, we use a **higher performance directive**, *omp atomic*

- *omp atomic* can only **protect** statements with a **single assignment**.
- For example, *omp atomic* protect statements of the following form:

```
x <op>= <expression>;
x++;
++x;
x--;
--x;
```

- < expression > **must not reference** *x*
 - < op > can be +, *, -, /, &, ^, |, <<, or >>
- *atomic* directive has special for **load – modify – store** instructions in modern processors, which is **more efficient** than *critical section*

Locks

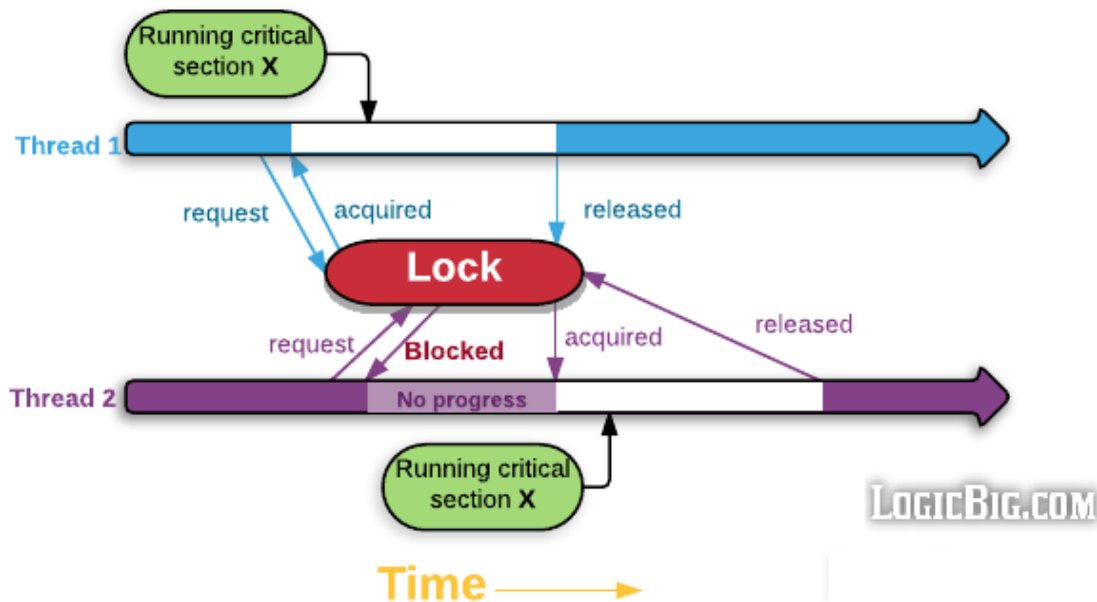
- As an alternative to *critical* sections, we can use **Locks**.
- A *lock* consists of a **data structure** and **functions** enforce mutual exclusion to data or resources.
- Defining locks follows this pseudocode:

```
/* Executed by one thread */
Initialize the lock data structure;
. . .
/* Executed by multiple threads */
Attempt to lock or set the lock data structure;
Critical section;
Unlock or unset the lock data structure;
. . .
/* Executed by one thread */
Destroy the lock data structure;
```

- The lock data structure is **shared among the threads** that will execute the critical section.
 - **One** of the threads will **initialize** the lock.
 - **One** of the threads will **destroy** the lock.
 - A thread entering the critical section will **set** the lock.
 - When the thread finishes the critical section, it will **releases** or **unset** the lock.

- How locks achieve mutual exclusion:

Mutual Exclusion of Critical Section



- Locks in OpenMP:

Action	Function
Initialize a lock	<code>void omp_init_lock (omp_lock_t* lock_p);</code>
Acquire a lock	<code>void omp_set_lock (omp_lock_t* lock_p);</code>
Release a lock	<code>void omp_unset_lock (omp_lock_t* lock_p);</code>
Remove a lock	<code>void omp_unset_lock (omp_lock_t* lock_p);</code>

Using locks in the message-passing program

- In the previous program, we used **critical directive** to enforce a mutually exclusive access to a **shared resource**.
 - `done_sending++;`
 - `Enqueue(q_p, my_rank, mesg);`
 - `Dequeue(q_p, &src, &mesg);`
- There is one efficiency issue with this critical section: the **critical** directive **blindly** allows only one thread to **Enqueue** or **Dequeue**.
- But we don't need to block threads that **don't conflict with each other**: for instance, it's safe for **thread 0** to enqueue a message in **thread 1's queue** at the same time that **thread 1** is enqueueing a message in **thread 2's queue**.
 - **So, the *critical* directive allows only one thread do enqueueing at a time regardless the process causes conflicts or not.**
- To **overcome** this, we use **locks** and instead of making the **whole function** of **enqueue** or **dequeue** a **critical section**, we will make the **corresponding queues only critical sections**.

- So, we can do the following:

```
# pragma omp critical
/* q_p = msg_queues[dest] */
Enqueue(q_p, my_rank, mesg);
```



```
/* q_p = msg_queues[dest] */
omp_set_lock(&q_p->lock);
Enqueue(q_p, my_rank, mesg);
omp_unset_lock(&q_p->lock);
```

```
# pragma omp critical
/* q_p = msg_queues[my_rank] */
Dequeue(q_p, &src, &mesg);
```



```
/* q_p = msg_queues[my_rank] */
omp_set_lock(&q_p->lock);
Dequeue(q_p, &src, &mesg);
omp_unset_lock(&q_p->lock);
```

- Check the source code in *queue_lk.h*, *queue_lk.c*, and *omp_msglk.c*.
 - The **lock** is defined as a data member in the **queue** struct.
- Now when a thread tries to send or receive a message, it can only be **blocked** by a thread accessing the **same message queue**, since **different message queues have different locks**.
- In previous implementation, **only one thread** could **send** at a time, regardless of the destination.

Critical directives, *atomic* directives, or locks?

- So far, we introduced **three** basic mutual exclusion mechanisms:

Mutual Exclusion

critical directive

atomic

locks

named

unnamed

Mechanism	Properties	Syntax	Group regions under one directive?
<i>atomic</i>	Simple and the fastest	#pragma omp atomic x <op>= <expression>; x++; --x;	Yes
<i>critical</i> (unnamed)	Easier than locks	#pragma omp critical { ... }	Yes
<i>critical</i> (named)	-	#pragma omp critical(name) { ... }	No
Locks	Better for data structures	omp_set_lock(lock); ... omp_unset_lock(lock);	No

- Critical regions specified by *atomic* or *critical* (unnamed) directives **may** treat all their regions as one block.
- For example:

```
#pragma omp parallel
{
    #pragma omp atomic
    var1 = var1 + 1; // Atomic operation on var1

    #pragma omp atomic
    var2 = var2 + 1; // Atomic operation on var2
}
```

- **In one implementation**, the runtime might enforce exclusive access, meaning that the operations on *var1* and *var2* cannot happen at the same time, even though they operate on different variables.
- **In another implementation**, the runtime might allow these operations to run concurrently if they access different variables, as there's no dependency or conflict between the two operations.
- The same applies if using **(unnamed)** *critical* directive:

```
#pragma omp parallel
{
    #pragma omp critical
    {
        var1 = var1 + 1;
    }

    #pragma omp critical
    {
        var2 = var2 + 1;
    }
}
```

- **Named *critical* directive and **locks**** avoid that.

```
#pragma omp parallel
{
    #pragma omp critical(section1)
    {
        var1 = var1 + 1;  // Atomic operation on var1
    }

    #pragma omp critical(section2)
    {
        var2 = var2 + 1;  // Atomic operation on var2
    }
}
```

Some caveats

1. **Don't mix** the different types of mutual exclusion for a **single** critical section.

- The code below uses two mutual exclusive mechanisms for the variable x .

# pragma omp atomic	# pragma omp critical
$x += f(y);$	$x = g(x);$


- It's possible that the two sections get executed **concurrently**, leading to **incorrect** results.
 - You can either: **use *critical* directive for both sections** or **rewrite $g()$ to have the form required by the *atomic* directive.**
2. There is no **guarantee of fairness** in mutual exclusion constructs.
 - For example, the code below may allow one thread to **always** be executing the function g , thus preventing other threads from accessing it.

```
while (1) {
    . . .
    # pragma omp critical
    x = g(my_rank);
    . . .
}
```

- This won't happen if the *while* loop **terminates**.
3. It can be dangerous to "**nest**" mutual exclusion constructs.

```
# pragma omp critical
y = f(x);
. . .
double f(double x) {
    # pragma omp critical
    z = g(x);  /* z is shared */
    . . .
}
```

- This will cause **deadlock**: a situation in which a group of threads are waiting for each other thread to finish.



- If a thread is executing the first block, it won't be able to **enter the second block**. At the same time, it **will not leave the first block** until it proceeds to the second block.
- One possible solution is use **named critical sections**

```

# pragma omp critical(one)
y = f(x);
. . .
double f(double x) {
#   pragma omp critical(two)
    z = g(x); /* z is global */
    . . .
}
  
```

- But this is not the **ultimate** solution, as deadlocks can occur as following:

Time	Thread <i>u</i>	Thread <i>v</i>
0	Enter crit. sect. one	Enter crit. sect. two
1	Attempt to enter two	Attempt to enter one
2	Block	Block

Caches, cache coherence, and false sharing