

# Introduction to Parallel Computing


Shared-memory programming  
with OpenMP

## Table of Contents

<i>Introduction</i> .....	2
<i>The trapezoidal rule</i> .....	6
Parallelizing the trapezoidal rule .....	9
<i>Scope of variables</i> .....	13

## Introduction

- OpenMP is an API for shared-memory Multiple Instruction Multiple Data (MIMD) programming.
- OpenMP is designed for systems in which each thread/process has access to all available memory.



**FIGURE 5.1**

A shared-memory system.

- It allows the programmer to **state that a block of code should be executed in parallel**, and the precise determination of the tasks and which thread should execute them is left to the **compiler and the run-time system**.
- Some compilers and languages that support OpenMP:
  - GNU: GCC – C/C++/Fortran
  - LLVM: Clang – C/C++
  - Microsoft: MSVC – C/C++
  - Nvidia: C/C++/Fortran
- OpenMP API is directive-based.
  - **Directives** are special preprocessor instructions called **pragma**.
- The **#pragma** directive is a special purpose directive that is used to turn on or off some features.

### Program 5.1: A “hello, world” program that uses OpenMP.

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void hello(void);

int main(int argc, char *argv[]) {
    int thread_count = strtol(argv[1], NULL, 10);

    #pragma omp parallel num_threads(thread_count)
    hello();

    return 0;
}

void hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);
}
```

- *strtol* is a function in the *stdlib.h* header that converts strings to long integer. It's used to convert the command line argument *argv[1]* into a string. The first parameter is string value to be converted, the second is a reference to the end of the string, the third is the base of long value.
- *omp parallel num\_threads()* tells the compiler that next line will be executed in parallel by a specific number of threads.
- *omp\_get\_thread\_num()*: a function returns the number (rank) of the current thread.
- *omp\_get\_num\_threads()*: a function returns the total number of threads that are assigned to the process.
- To compile and run the previous program:
  - For Clang and Apple M1

```
clang -Xpreprocessor -fopenmp -I/opt/homebrew/opt/libomp/include -
L/opt/homebrew/opt/libomp/lib -lomp -o my_program <filename.c>
```

- For GCC on Windows, Linux, and Mac

```
gcc-14 -o my_program -fopenmp <filename.c>
```

- Or

```
gcc -o my_program -fopenmp <filename.c>
```

- Then, run using

```
./my_program 4
```

- The program output is different from a run to another. This is because each thread is competing for access to *stdout*.


```
Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 3 of 4

Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4
Hello from thread 3 of 4
```


or

```
Hello from thread 3 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4
```

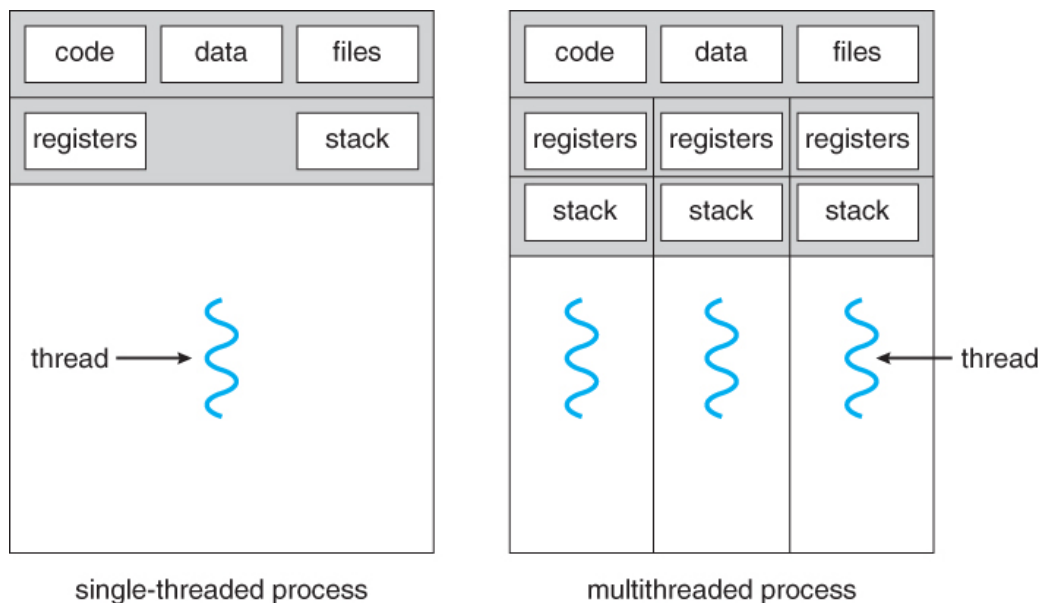
- When the program **reaches the *parallel* directive**, the original thread **forks** *thread\_count* – 1 additional threads.
- The collection of threads executing the *parallel* block – the original thread and the new threads – is called a **team**.
  - **Master** thread: the first thread of execution, thread 0
  - **Parent** thread: thread that encountered a *parallel* directive and started a team of threads.
    - In many cases, the parent is also the master thread.
  - **Child** thread: each thread started by the parent.



- **Implicit barrier:** a thread that has completed the parallel block of code will wait for all the other threads in the team to complete the block.



- Each thread has its own **stack**; so, a thread executing the *Hello* function will create its own private, local variables in the function.
- Since ***stdout*** is shared among the threads, each thread can execute the *printf()* to print its rand and the number of threads.
  - There is no **scheduling** of access to *stdout*, so the actual order in which the threads print their results is **nondeterministic**.



- To handle missing *omp* or if a compiler doesn't support *omp*, we can modify the previous program as follows:

```
#include <stdio.h>
#include <stdlib.h>

#ifdef _OPENMP
#include <omp.h>
#endif

void hello(void);

int main(int argc, char *argv[]) {
    int thread_count = strtol(argv[1], NULL, 10);

    #pragma omp parallel num_threads(thread_count)
    hello();


    return 0;
}

void hello(void) {
#ifdef _OPENMP
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
#else
    int my_rank = 0;
    int thread_count = 1;
#endif


    printf("Hello from thread %d of %d\n", my_rank, thread_count);
}
```

## The trapezoidal rule

- The trapezoidal rule is used to **estimate the area under the curve**.
- To estimate the area between the graph  $f(x)$ , the vertical lines  $x = a$  and  $x = b$  and the x-axis, divide the interval  $[a, b]$  into  $n$  subintervals and approximating the area over each subinterval by the area of a trapezoid.



- The area of each subinterval is calculated using:  $\frac{h}{2}(f(x_i) + f(x_{i+1}))$



- $h = x_{i+1} - x_i \rightarrow$  the length of each subinterval
  - $f(x_i)$  and  $f(x_{i+1})$  is the length of the vertical segments.
- Since we are dividing the area into subintervals, we let
  - $n$  is number of subintervals
  - $b$  and  $a$  are the limits of the intervals
  - $h = (b - a)/n$ , the length of each subinterval
  - $x_i = a + ih$  for  $i = 0, 1, \dots, n$
- Thus, if we call the leftmost endpoint  $x_0$ , and the rightmost endpoint  $x_n$ , we have that:  
 $x_0 = a, \quad x_1 = a + h, \quad x_2 = a + 2h, \dots, \quad x_{n-1} = a + (n - 1)h, \quad x_n = b$
- The sum of the areas =  $h \left( \frac{f(x_0)}{2} + f(x_1) + f(x_2) + \dots + f(x_{n-1}) + \frac{f(x_n)}{2} \right)$
- Then the approximation is

```

/* Input:  a, b, n */
h = (b-a)/n;
approx = (f(a) + f(b))/2.0;
for (i = 1; i <= n-1; i++) {
    x_i = a + i*h;
    approx += f(x_i);
}
approx = h*approx;

```

```

#include <stdio.h>

double f(double x);    /* Function we're integrating */
double Trap(double a, double b, int n, double h);

int main(void) {
    double integral;    /* Store result in integral */
    double a, b;        /* Left and right endpoints */
    int n;              /* Number of trapezoids */
    double h;

    printf("Enter a, b, and n\n");
    scanf("%lf", &a);
    scanf("%lf", &b);
    scanf("%d", &n);

    h = (b - a) / n;
    integral = Trap(a, b, n, h);

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.15f\n", a, b, integral);

    return 0;
} /* main */

double Trap(double a, double b, int n, double h) {
    double integral = (f(a) + f(b)) / 2.0;
    for (int k = 1; k <= n - 1; k++) {
        integral += f(a + k * h);
    }
    integral = integral * h;

    return integral;
} /* Trap */

double f(double x) {
    return x * x;
} /* f */

```



## Parallelizing the trapezoidal rule

### Partition the problem solution into tasks

- Find the area of a single trapezoid
- Computing the sum of these areas

### Identify communication channels


- Join the task of computing the area of each trapezoid with the task of computing the final sum

### Aggregate tasks into composite tasks

- Assign a contiguous block of trapezoids to each thread.
- Partition the interval  $[a, b]$  into smaller subintervals.


### Map composite tasks to cores

- Each thread applies the serial trapezoidal rule to its subinterval
- Each thread will be executed on a single core



**FIGURE 3.5**

Tasks and communications for the trapezoidal rule.



Parallelizing the problem:

- Assume we have  $a = 0$ ,  $b = 500$ ,  $n = 100$  and the number of threads = 4

Compute the length of each subinterval

- $h = (b - a)/n \Rightarrow h = (500 - 0)/100 = 5$

Compute the number of intervals assigned to each thread

- $local_n = n/num\_threads \Rightarrow local_n = 100/4 = 25$

Compute the start and the end of sub-areas assigned to each thread

$$local_a = a + rank * local_n * h \quad local_b = local_a + local_n * h$$

- $local_a = 0 + 0 * 25 * 5 = 0$  ,  $local_b = 0 + 25 * 5 = 125$
- $local_a = 0 + 1 * 25 * 5 = 125$  ,  $local_b = 125 + 25 * 5 = 250$
- $local_a = 0 + 2 * 25 * 5 = 250$  ,  $local_b = 250 + 25 * 5 = 375$
- $local_a = 0 + 3 * 25 * 5 = 375$  ,  $local_b = 375 + 25 * 5 = 500$

[0:125]

[125: 250]

[250:375]

[375:500]

$(f(0) + f(125)) / 2$

$(f(125) + f(250)) / 2$

$(f(250) + f(375)) / 2$

$(f(375) + f(500)) / 2$

```
for (i = 1; i <= local_n - 1; i++) {  
    x = local_a + i * h;  
    my_result += f(x);  
}  
my_result = my_result * h;
```

Final result = my\_result + my\_result + my\_result + my\_result

```

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>

double f(double x); /* Function we're integrating */
void trap(double a, double b, int n, double *global_result_p);

int main(int argc, char *argv[]) {
    double global_result = 0.0; /* Store result in global_result */
    double a, b; /* Left and right endpoints */
    int n; /* Total number of trapezoids */

    int thread_count = strtol(argv[1], NULL, 10);
    printf("Enter a, b, and n\n");
    scanf("%lf %lf %d", &a, &b, &n);

    # pragma omp parallel num_threads(thread_count)
    trap(a, b, n, &global_result);

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.14e\n", a, b, global_result);
    return 0;
} /* main */

double f(double x) {
    return x * x;
} /* f */

void trap(double a, double b, int n, double *global_result_p) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
    double x;

    double h = (b - a) / n;
    int local_n = n / thread_count;
    double local_a = a + my_rank * local_n * h;
    double local_b = local_a + local_n * h;
    double my_result = (f(local_a) + f(local_b)) / 2.0;
    for (int i = 1; i <= local_n - 1; i++) {
        x = local_a + i * h;
        my_result += f(x);
    }
    my_result = my_result * h;

    # pragma omp critical
    *global_result_p += my_result;
} /* trap */

```

- The program reads the number of threads to fork as a command line argument.
- Next, the program asks the user to enter: ***a***, the start of the interval, ***b***, the end of the interval, and ***n*** the number of subintervals (the smaller trapezoids).

- Notice that when we call the *trap* function, we pass *a, b, n* as values, while the *global\_result* is passed as an address.
- Inside *trap()*:
  - The program gets the **rank** (ID) of the current thread. This used to compute the start and the end of the interval that is assigned to the current thread.
  - *thread\_count* is used to store the total number of threads used to execute the function. This is used to compute how many subintervals (smaller trapezoids) to each thread.
  - *h* is the **length** value of each trapezoid, which is computed by dividing the size of the interval ( $b - a$ ) by the total number of trapezoids ( $n$ ).
  - *local\_n* is a variable that computes how many trapezoids are assigned to each thread. It is computed by dividing the total number of trapezoids by the total number of threads available.
  - *local\_a* and *local\_b* store the limits of each subinterval (trapezoid) assigned to each thread.
  - *my\_result* is **local integral value** of each thread. It's computed using the *local\_a* and *local\_b* instead of the limits of the entire interval  $[a, b]$ .
  - *#pragma omp critical* is a directive that tells the compiler to make the next statement to have **mutually exclusive** access. This means that the statement  $*global\_result += my\_result$  will be executed by only one thread. If more than one thread is executing the same statement at the same time, the result will be incorrect.
  - A code that includes access to a shared resource and causes a **race condition**, is called a **critical section**.
    - **Race condition:** multiple threads are attempting to access a shared resource, at least one of the accesses is an update, and the accesses can result in an error.
- Notice that the variables *local\_a*, *local\_b*, and *my\_result* are different for each thread.
- If number of total trapezoids,  $n$ , is **not divisible** by the number of available threads, *thread\_count*, then we will result in **inaccurate** results and we must include error handling mechanisms.
- For example, if  $n = 14$  and *thread\_count* = 4, then each thread will compute
 
$$local\_n = n / thread\_count = 14 / 4 = 3$$
 Thus, each thread will use 3 trapezoids and *global\_result* will be computed with  $4 \times 3 = 12$  instead of 14.
- We can handle this error by including an *if* condition to check that  $n$  is divisible by *thread\_count*.

```
if ( n % thread_count != 0 ) {
    fprintf (stderr, "n must be evenly divisible by thread_count\n" );
    exit (0) ;
}
```

- During the execution of the program, the  $local\_a = a + my\_rank * local\_n * h$ ; will be as following:

```


thread 0:  a + 0*local_n*h
thread 1:  a + 1*local_n*h
thread 2:  a + 2*local_n*h
. . .

```

- So, the value of  $my\_rank$  is replaced by the rank of the thread.
- The same applies for the  $local\_b = local\_a + local\_n * h$ ;


## Scope of variables

- In serial programming, the **scope** of a variable consists of those parts of a program in which the variable can be used.
- In OpenMP, the **scope** of a variable refers to the set of threads that can access the variable in a parallel block.



### Shared variable

all team threads can access it



### Private variable

Only one thread can access it

- In the *trapezoidal* program:
  - The variables declared in the **main function** ( $a, b, n, global\_result, thread\_count$ ) are **shared**; they are declared before the *parallel* directive.
  - The variables inside the **trap** function ( $h, x, my\_rank$ , etc.) are **private**; they are declared within the scope of parallel directive.
  - The variable  **$global\_result\_p$**  in the *trap* function is **private**. However, it refers to the **shared** variable  **$global\_result$**  that is defined in the main function. So,  **$global\_result\_p$**  is treated as a **shared variable**.
  - If  **$global\_result\_p$**  were **private** to each thread, there would be no need for the *critical* directive. It must be shared to compute the final result computed by individual threads.

- In summary, variables that have been declared before a *parallel* directive have **shared** scope, while variables declared in the block (e.g., local variables in functions) have **private** scope.