

# Introduction to Parallel Computing

Shared-memory programming  
with OpenMP

## Table of Contents

<i>Scheduling loops</i> .....	2
-------------------------------	---

## Scheduling loops

- The *parallel for* directive automatically applies **block partitioning**: if the serial loop has  $n$  iterations, then thread 0 handles the first  $n/\text{thread\_count}$  iterations, thread 1 handles the second  $n/\text{thread\_count}$  iterations, and so on.
- This could be **less optimal**.
- For example, say we want to parallelize the serial loop

```
sum = 0.0;
for (i = 0; i <= n; i++)
    sum += f(i);
```

Where the time required to compute  $f$  is **proportional** to the size of the argument  $i$ .

- Then, block partitioning will assign much more work to the last thread,  $\text{thread\_count} - 1$ , than it will assign to thread 0
- It's better to use **cyclic partitioning** for better **work-balance** between threads: the iterations are assigned, one at a time, in a “round-robin” fashion to the threads.
- Suppose  $t = \text{thread\_count}$ . Then a **cyclic partitioning** will assign the iterations as follows

Thread	Iterations
0	0, $n/t$ , $2n/t$ , ...
1	1, $n/t + 1$ , $2n/t + 1$ , ...
$\vdots$	$\vdots$
$t - 1$	$t - 1$ , $n/t + t - 1$ , $2n/t + t - 1$ , ...

- Example:

```
double f(int i) {
    int j, start = i*(i+1)/2, finish = start + i;
    double return_val = 0.0;

    for (j = start; j <= finish; j++) {
        return_val += sin(j);
    }
    return return_val;
} /* f */
```

- Run the next three programs and compare the performance.  
`gcc-14 -o prog -fopenmp <filename.c>; ./prog`

- The serial program

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>

double Sum(long n);
double f(long i);

int main(int argc, char* argv[]) {
    double global_result;
    double start, finish;
    long n = 100000;

    start = omp_get_wtime();
    global_result = Sum(n);
    finish = omp_get_wtime();

    printf("Result = %.14e\n", global_result);
    printf("Elapsed time = %f seconds\n", finish-start);

    return 0;
}

double f(long i) {
    long j;
    long start = i*(i+1)/2;
    long finish = start + i;
    double return_val = 0.0;

    for (j = start; j <= finish; j++) {
        return_val += sin(j);
    }
    return return_val;
}

double Sum(long n) {
    double approx = 0.0;
    long i;

    for (i = 0; i <= n; i++) {
        approx += f(i);
    }
    return approx;
}
```

- The parallel program

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>

double Sum(long n, int thread_count);
double f(long i);

int main(int argc, char* argv[]) {
    double global_result;
    double start, finish;

    int thread_count = 2;
    long n = 100000;

    start = omp_get_wtime();
    global_result = Sum(n, thread_count);
    finish = omp_get_wtime();

    printf("Result = %.14e\n", global_result);
    printf("Elapsed time = %f seconds\n", finish-start);

    return 0;
}

double f(long i) {
    long j;
    long start = i*(i+1)/2;
    long finish = start + i;
    double return_val = 0.0;

    for (j = start; j <= finish; j++) {
        return_val += sin(j);
    }
    return return_val;
}

double Sum(long n, int thread_count) {
    double approx = 0.0;
    long i;

    # pragma omp parallel for num_threads(thread_count) \
    reduction(+: approx)
    for (i = 0; i <= n; i++) {
        approx += f(i);
    }
    return approx;
}
```

- The **scheduled** parallel program

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>

double Sum(long n, int thread_count);
double f(long i);

int main(int argc, char *argv[]) {
    double global_result;
    double start, finish;

    int thread_count = 2;
    long n = 100000;

    start = omp_get_wtime();
    global_result = Sum(n, thread_count);
    finish = omp_get_wtime();

    printf("Result = %.14e\n", global_result);
    printf("Elapsed time = %f seconds\n", finish - start);

    return 0;
}

double f(long i) {
    long j;
    long start = i * (i + 1) / 2;
    long finish = start + i;
    double return_val = 0.0;

    for (j = start; j <= finish; j++) {
        return_val += sin(j);
    }
    return return_val;
}

double Sum(long n, int thread_count) {
    double approx = 0.0;
    long i;

    # pragma omp parallel for num_threads(thread_count) \
    reduction(+: approx) schedule(dynamic)
    for (i = 0; i <= n; i++) {
        approx += f(i);
    }
    return approx;
}
```

- **Scheduling** in OpenMP: assigning iterations to threads.
- The ***schedule* clause** can be used to assign iterations to threads:  
`schedule(< type > [, chunk size])`
- **Chunk**: a block of iterations that would be executed consecutively in the serial loop.
  - *chunk size* determines the **number of iterations in a block**. It's **optional**.

- Types of scheduling:

#### static

- The iterations are assigned to the threads before the loop is executed

#### dynamic/guided

- The iterations are assigned to the threads while the loop is executing
- After a thread completes its current set of iterations, it requests more from the run-time system

#### auto

- The compiler and/or the run-time system determine the schedule
- *Does not specify chunk size*

#### runtime

- The schedule is determined at run-time based on an environment variable

### The *static* schedule type

- The system assigns **chunks** of ***chunksize* iterations** to each thread in a round-robin fashion.
- Example: suppose we have 12 iterations, 0, 1,...,11, and three threads.
  - For *schedule(static, 1)*, the iterations will be assigned as

```
Thread 0 : 0, 3, 6, 9
Thread 1 : 1, 4, 7, 10
Thread 2 : 2, 5, 8, 11
```

- For *schedule(static, 2)*, the iterations will be assigned as

```
Thread 0 : 0, 1, 6, 7
Thread 1 : 2, 3, 8, 9
Thread 2 : 4, 5, 10, 11
```

- For *schedule(static, 4)*, the iterations will be assigned as

```
Thread 0 : 0, 1, 2, 3
Thread 1 : 4, 5, 6, 7
Thread 2 : 8, 9, 10, 11
```

- The **default** schedule is defined by a particular implementation of OpenMP, but in most cases it is equivalent to the clause:

*schedule( **static** , total\_iterations / thread\_count )*

- The *static* schedule is a good choice when each loop iteration takes roughly the **same amount of time** to compute.

### The *dynamic* and *guided* schedule type


- *dynamic* schedule: the iterations are also **broken up into chunks** of *chunksize* consecutive iterations.
  - Each thread executes a chunk, and when a thread finishes a chunk, it requests another one from the run-time system.
- The **default *chunksize*** is 1.
- The **difference** between *static* and *dynamic* schedules:
  - *dynamic* schedule assigns ranges to threads on a **first-come, first-served** basis.
    - Good choice if loop iterations **do not take a uniform** amount of time to compute.
    - There is some **overhead** associated with assigning them dynamically at run-time.
  - *static* schedule assigns a range of threads **fixedly** at the start of the loop
    - Used when loop iterations **take a uniform** amount of time.
- With **larger** chunk sizes, **fewer** dynamic assignments will be made.
- In a ***guided* schedule**, as chunks are completed, the size of the new chunks **decreases**.
- If no *chunksize* is specified, the size of the chunks decreases down to **1**. If *chunksize* is specified, it decreases down to ***chunksize***.

### The *runtime* schedule type


- When *schedule(runtime)* is specified, the system uses the **environment variable OMP\_SCHEDULE** to determine at run-time how to schedule the loop.
  - Environment variables are **named values** that can be accessed by a running program.
  - Examples: *PATH, HOME, SHELL*
- The **OMP\_SCHEDULE** environment variable can take on any of the values that can be used for a static, dynamic, or guided schedule.
  - `$ export OMP_SCHEDULE="static ,1"`

- Scheduling **visualization** for the **static**, **dynamic**, and **guided** schedule types with **4 threads and 32 iterations**.


`schedule(static)`




`schedule(static, 2)`



`schedule(dynamic, 2)`



`schedule(guided)`





### Which schedule type to choose?

- It depends on the problem and the nature of the loop.
- There is some **overhead** associated with the use of a **schedule clause**.
  - *static* overhead **less than** *dynamic* overhead **less than** *guided* overhead
- Practically, you should **try different options** until you hit a satisfactory performance.
  - Make use of the *schedule(runtime)* to test different settings.