

Introduction to Parallel Computing


Shared-memory programming with OpenMP

Table of Contents


<i>Producers and consumers</i>	2
Sending messages	3
Receiving messages	3
Terminating the program	4
Experiments	7
Locks	8
Using locks in the message-passing program	9
<i>Critical</i> directives, <i>atomic</i> directives, or locks?	10
Some caveats	12

Producers and consumers

- A **queue** is a data structure in which elements are inserted at the rear of the queue and removed from the front of the queue.
 - Adding elements → enqueue
 - Removing elements → dequeue



- Queues appear in many **applications**:
 - Print spooling
 - Operating systems scheduling/task scheduling
 - Data buffers in networking
 - Load balancing in web servers
- Queues are fundamental in **producer-consumer** applications:
 - Part the of the application is **producing** data, and another part is **consuming** the data



- Let's implement a **message-passing** application that is based on the **producer-consumer** pattern.
 - Each thread has a **shared-message queue**.
 - Each thread generates random (integer values) messages and random destination for the messages.
 - A **producer** thread **enqueues** the message in the queue.
 - A **consumer** thread checks if the queue received a message, and then **dequeues** it.
 - Each thread **alternates** between sending and receiving messages.

- The user will specify the **number of messages** each thread should send.
 - When a thread is **done sending messages**, it receives messages until all the threads are done, at which point all the threads **quit**.
- Pseudocode for each thread might look something like this:

```
for (sent_msgs = 0; sent_msgs < send_max; sent_msgs++) {
    Send_msg();
    Try_receive();
}

while (!Done())
    Try_receive();
```

- Core segments/functions of the program:
 - Sending messages → *Send_msg()*
 - Receiving messages → *Try_receive()*
 - Terminating the program → *Done()*

Sending messages

- Accessing a **message queue** to **enqueue** a message is a **critical section**.
 - We need to regulate how threads access the queues to avoid dropping messages.
- To successfully **enqueue** a message, we need a pointer to the **rear** of the queue.
 - When we enqueue a message, we'll update the rear pointer.

```
msg = random();
dest = random() % thread_count;
# pragma omp critical
    Enqueue(queue, dest, my_rank, msg);
```

- The *Enqueue* functions handles updating the rear pointer.
 - Note that this allows a thread to send a message to itself.

Receiving messages

- **Only the owner of the queue (that is, the destination thread) will dequeue from a given message queue.**
- If there are at least two messages in the queue, a call to *Dequeue* **can't conflict** with any calls to *Enqueue*.
- Before *Dequeuing* the queue, **we must check its size**; if 0, return none. If the size is ≥ 1 , then return the element at the front.
 - We store two variables: *enqueued*, which is the number of added elements, and *dequeued*, which is the number of removed elements.

```

queue_size = enqueued - dequeued;
if (queue_size == 0) return;
else if (queue_size == 1)
#   pragma omp critical
    Dequeue(queue, &src, &mesg);
else
    Dequeue(queue, &src, &mesg);
    Print_message(src, mesg);

```

- The *critical section* is added to prevent other threads from updating the queue when the owner thread is *Dequeueing* it when it has only one element.

Terminating the program

- The **naïve way** to terminate the program (implementing *Done()*) is as follows:

```

queue_size = enqueued - dequeued;
if (queue_size == 0)
    return TRUE;
else
    return FALSE;

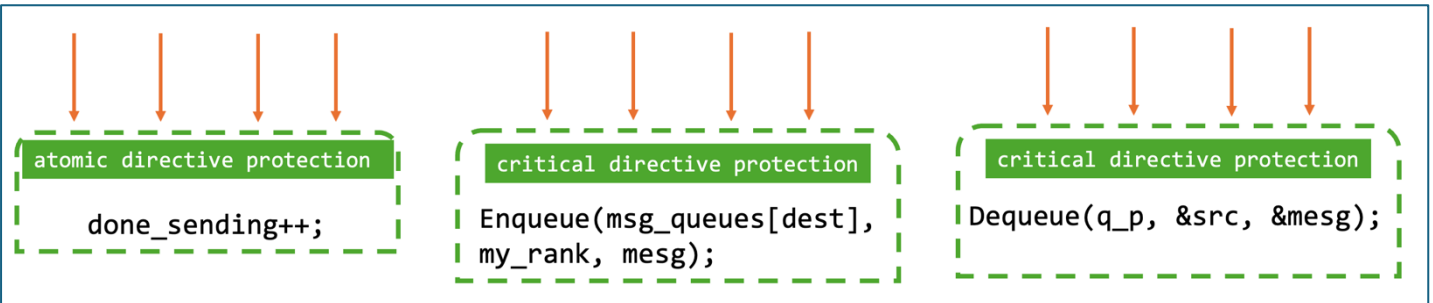
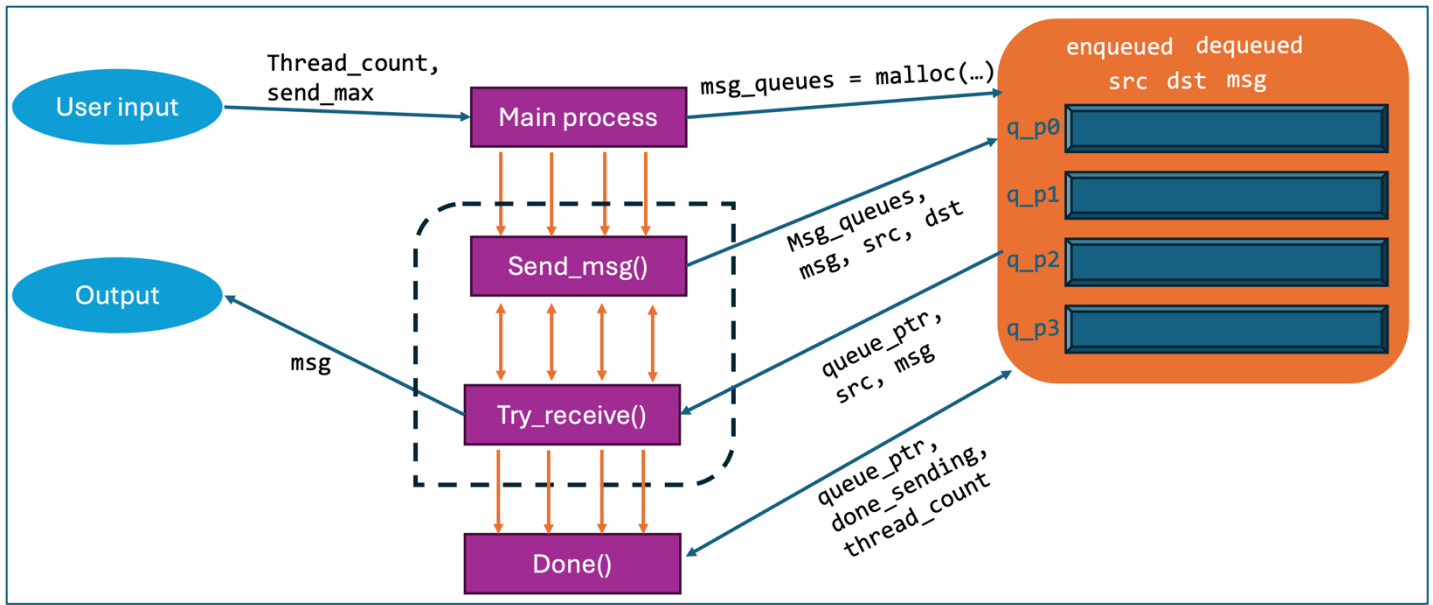
```

- The problem with this way is that it is possible that the owner thread compute *queue_size* = 0, while another thread is adding an element into the queue.
 - Thus, the added message will never be received by the owner thread.
- Instead, we can include *done_sending* as a counter variable, so it serves as a flag when all threads have no more messages to send.

```

queue_size = enqueued - dequeued;
if (queue_size == 0 && done_sending == thread_count)
    return TRUE;
else
    return FALSE;

```



```

#include <stdio.h>
#include <omp.h>
#include <stdlib.h>
#include "queue.h"
// gcc-14 -o prog -fopenmp -DDEBUG omp_msgps.c queue.c

/*-----*/
void Send_msg(struct queue_s *msg_queues[], int my_rank, int
thread_count, int msg_number) {
    int mesg = -msg_number;
    int dest = random() % thread_count;
    # pragma omp critical
    Enqueue(msg_queues[dest], my_rank, mesg);
    printf("Thread %d > sent %d to %d\n", my_rank, mesg, dest);
}
/*-----*/
void Try_receive(struct queue_s *q_p, int my_rank) {
    int src, mesg;
    int queue_size = q_p->enqueued - q_p->dequeued;

    if (queue_size == 0) return;
    else if (queue_size == 1)
    # pragma omp critical
        Dequeue(q_p, &src, &mesg);
    else
        Dequeue(q_p, &src, &mesg);
    //printf("Thread %d > received %d from %d\n", my_rank, mesg, src);
}
  
```

```

/*-----*/
int Done(struct queue_s *q_p, int done_sending, int thread_count) {
    int queue_size = q_p->enqueued - q_p->dequeued;
    if (queue_size == 0 && done_sending == thread_count)
        return 1;
    else
        return 0;
} /* Done */
/*-----*/

int main(int argc, char *argv[]) {
    int thread_count;
    int send_max;
    struct queue_s **msg_queues;
    int done_sending = 0;

    thread_count = 4;
    send_max = 5;

    msg_queues = malloc(thread_count * sizeof(struct queue_node_s *));

    # pragma omp parallel num_threads(thread_count) \
        default(none) shared(thread_count, send_max, msg_queues,
done_sending)
    {
        int my_rank = omp_get_thread_num();
        int msg_number;
        srand(my_rank);
        msg_queues[my_rank] = Allocate_queue();

    # pragma omp barrier /* Don't let any threads send messages */
        /* until all queues are constructed */

        for (msg_number = 0; msg_number < send_max; msg_number++) {
            Send_msg(msg_queues, my_rank, thread_count, msg_number);
            Try_receive(msg_queues[my_rank], my_rank);
        }

        // done_sending is a critical section
        // using atomic instead of critical section has a better performance
    # pragma omp atomic
        done_sending++;

        while (!Done(msg_queues[my_rank], done_sending, thread_count))
            Try_receive(msg_queues[my_rank], my_rank);

        /* My queue is empty, and everyone is done sending */
        /* So my queue won't be accessed again, and it's OK to free it */
        Free_queue(msg_queues[my_rank]);
        free(msg_queues[my_rank]);
    } /* omp parallel */

    free(msg_queues);
    return 0;
} /* main */

```

- When the program begins execution, the master thread will get command-line arguments and allocate an array of message queues, one for each thread.
- This array needs to be shared among the threads, since any thread can send to any other thread, and hence any thread can enqueue a message in any of the queues.
- We use *omp barrier* to add an explicit barrier so that no thread starts enqueueing or dequeuing until all threads have constructed their queues to avoid errors.
- After completing its sends, each thread increments *done_sending* before proceeding to its final loop of receives.
 - Incrementing *done_sending* is a **critical section**, and we could protect it with a **critical directive**.
 - Instead, we use a **higher performance directive**, *omp atomic*
- *omp atomic* can only **protect** statements with a **single assignment**.
- For example, *omp atomic* protect statements of the following form:

```
x <op>= <expression>;
x++;
++x;
x--;
--x;
```

- *< expression >* **must not reference** *x*
 - *< op >* can be +, *, -, /, &, ^, |, <<, or >>
- *atomic* directive has special for **load – modify – store** instructions in modern processors, which is **more efficient** than *critical section*

Experiments

1. Run the program as is – with **4 threads and 5 messages** to send and check the output.
 - a. You will see that each thread must **send exactly 5 messages**.
 - b. Threads **may or may not** receive 5 messages.
2. **Remove the critical section** from the *Try_receive* function and run the program **with 4 threads and 5 messages**. The *Try_receive* function should look like this:

```
/*-----*/
void Try_receive(struct queue_s *q_p, int my_rank) {
    int src, mesg;
    int queue_size = q_p->enqueued - q_p->dequeued;
    if (queue_size == 0) return;
    Dequeue(q_p, &src, &mesg);
    //printf("Thread %d > received %d from %d\n", my_rank, mesg, src);
} /* Try_receive */
```

- a. The program works fine without any issues.

3. **Revert the *Try_receive* function** and run the program with **10 threads and 10000** messages.
 - a. You notice that the program may **take a little longer, but it finishes successfully.**
4. **Remove the critical section** again from *Try_receive* and run the program with **10 threads and 10000 messages.**
 - a. **It will either run for ever or it will crash because of heap errors.**

```

Thread 2 > sent -1054 to 3
prog(11076,0x16b79b000) malloc: Heap corruption detected, free list is damaged at 0x60000004ebf0
*** Incorrect guard value: 0
Thread 5 > sent -1063 to 4
Thread 4 > sent -1012 to 5
Thread 3 > sent -1097 to 1
prog(11076,0x16bdbf000) malloc: *** error for object 0x600000004c3e0: pointer being freed was not allocated
prog(11076,0x16bdbf000) malloc: *** set a breakpoint in malloc_error_break to debug
Thread 0 > sent -1149 to 5
prog(11076,0x1eb99c840) malloc: *** error for object 0x600000004c350: pointer being freed was not allocated
przsh: abort      ./prog

```

5. Run the **same program** but with **100 threads and 100,000 messages.**
 - a. The program is **very likely** to run for ever.
 6. **Revert the *Try_recieve* function** and run with **100 threads and 100000** messages.
 - a. The program **takes a few minutes** and **finishes successfully.**
- *These experiments demonstrate the importance of handling **special edge-cases**, in which the termination of a thread/program is dependent on state of the object (whether the queue is empty or not). With small inputs (e.g., 4 threads and 5 messages) errors **may not** occur. But when the problem size's increases (10 threads and 100000 messages) problems occur and hard to detect.*
 - Go through the same program but with the textbook version.

Locks

- As an alternative to *critical* sections, we can use **Locks**.
- A *lock* consists of a **data structure/functions** enforce mutual exclusion to data or resources.
- Defining locks follows this pseudocode:

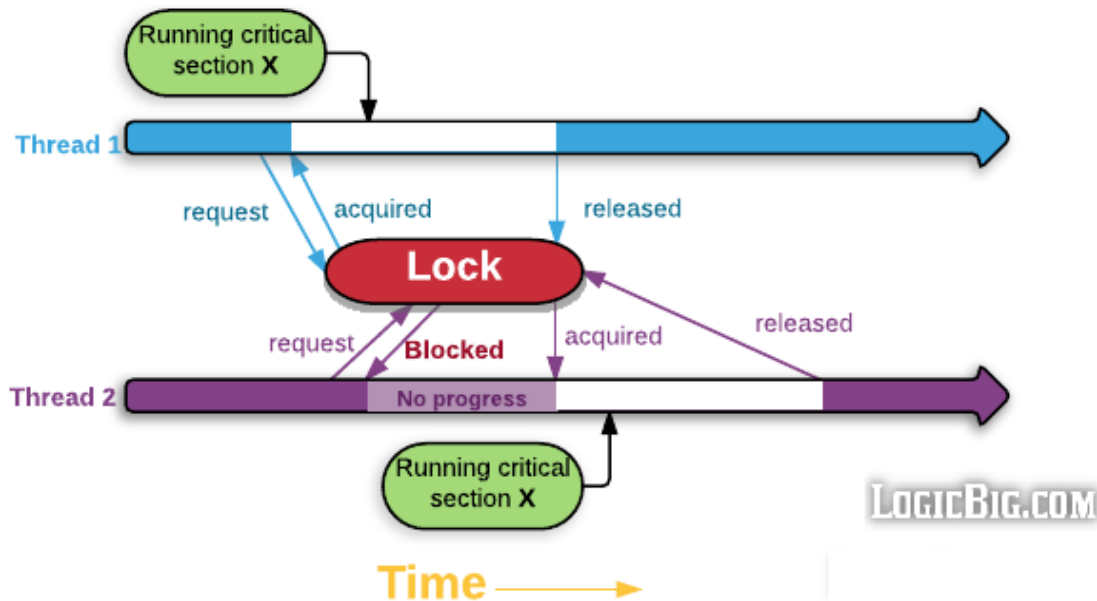
```

/* Executed by one thread */
Initialize the lock data structure;
. . .
/* Executed by multiple threads */
Attempt to lock or set the lock data structure;
Critical section;
Unlock or unset the lock data structure;
. . .
/* Executed by one thread */
Destroy the lock data structure;

```


- The lock data structure is **shared among the threads** that will execute the critical section.
- **One** of the threads will **initialize** the lock.
- **One** of the threads will **destroy** the lock.
- A thread entering the critical section will **set** the lock.
- When the thread finishes the critical section, it will **releases** or **unset** the lock.
- How locks achieve mutual exclusion:

Mutual Exclusion of Critical Section



- Locks in OpenMP:

Action	Function
Initialize a lock	<code>void omp_init_lock (omp_lock_t* lock_p);</code>
Acquire a lock	<code>void omp_set_lock (omp_lock_t* lock_p);</code>
Release a lock	<code>void omp_unset_lock (omp_lock_t* lock_p);</code>
Remove a lock	<code>void omp_destroy_lock (omp_lock_t* lock_p);</code>

Using locks in the message-passing program

- In the previous program, we used **critical directive** to enforce a mutually exclusive access to a **shared resource**.
 - `done_sending++;`
 - `Enqueue(q_p, my_rank, mesg);`
 - `Dequeue(q_p, &src, &mesg);`
- There is one efficiency issue with this critical section: the *critical* directive **blindly** allows only one thread to **Enqueue** or **Dequeue**.

- But we don't need to block threads that **don't conflict with each other**: for instance, it's safe for **thread 0** to enqueue a message in **thread 1's queue** at the same time that **thread 1** is enqueueing a message in **thread 2's queue**.
 - So, the **critical** directive allows only one thread do enqueueing at a time regardless the process causes conflicts or not.
- To **overcome** this, we use *locks* and instead of making the **whole function** of *enqueue* or *dequeue* a **critical section**, we will make the **corresponding queues only critical sections**.
- So, we can do the following:

```
# pragma omp critical
/* q_p = msg_queues[dest] */
Enqueue(q_p, my_rank, mesg);
```



```
/* q_p = msg_queues[dest] */
omp_set_lock(&q_p->lock);
Enqueue(q_p, my_rank, mesg);
omp_unset_lock(&q_p->lock);
```

```
# pragma omp critical
/* q_p = msg_queues[my_rank] */
Dequeue(q_p, &src, &mesg);
```

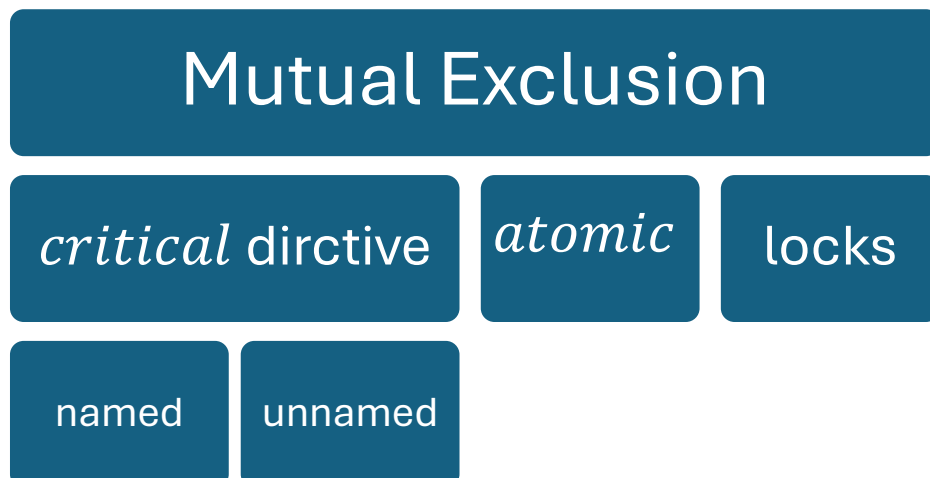


```
/* q_p = msg_queues[my_rank] */
omp_set_lock(&q_p->lock);
Dequeue(q_p, &src, &mesg);
omp_unset_lock(&q_p->lock);
```

- Check the source code in *queue_lk.h*, *queue_lk.c*, and *omp_msglk.c*.
 - The **lock** is defined as a data member in the *queue* struct.
- Now when a thread tries to send or receive a message, it can only be **blocked** by a thread accessing the **same message queue**, since **different message queues have different locks**.
- In previous implementation, **only one thread** could **send** at a time, regardless of the destination.

Critical directives, atomic directives, or locks?

- So far, we introduced **three** basic mutual exclusion mechanisms:



Mechanism	Properties	Syntax	Group regions under one directive?
<i>atomic</i>	Simple and the fastest	#pragma omp atomic x <op>= <expression>; x++; --x;	Yes
<i>critical</i> (unnamed)	Easier than locks	#pragma omp critical { ... }	Yes
<i>critical</i> (named)	-	#pragma omp critical(name) { ... }	No
Locks	Better for data structures	omp_set_lock(lock); ... omp_unset_lock(lock);	No

- Critical regions specified by *atomic* or *critical* (unnamed) directives **may** treat all their regions as one block.
- For example:

```
#pragma omp parallel
{
    #pragma omp atomic
    var1 = var1 + 1; // Atomic operation on var1

    #pragma omp atomic
    var2 = var2 + 1; // Atomic operation on var2
}
```

- **In one implementation**, the runtime might enforce exclusive access, meaning that the operations on *var1* and *var2* cannot happen at the same time, even though they operate on different variables.
- **In another implementation**, the runtime might allow these operations to run concurrently if they access different variables, as there's no dependency or conflict between the two operations.
- The same applies if using **(unnamed)** *critical* directive:

```
#pragma omp parallel
{
    #pragma omp critical
    {
        var1 = var1 + 1;
    }

    #pragma omp critical
    {
        var2 = var2 + 1;
    }
}
```

- **Named *critical* directive** and **locks** avoid that.

```
#pragma omp parallel
{
    #pragma omp critical(section1)
    {
        var1 = var1 + 1; // Atomic operation on var1
    }

    #pragma omp critical(section2)
    {
        var2 = var2 + 1; // Atomic operation on var2
    }
}
```

Some caveats

1. **Don't mix** the different types of mutual exclusion for a **single** critical section.
 - The code below uses two mutual exclusive mechanisms for the variable *x*.

# pragma omp atomic x += f(y);	# pragma omp critical x = g(x);
--	---


- It's possible that the two sections get executed **concurrently**, leading to **incorrect** results.
 - You can either: **use *critical* directive for both sections** or **rewrite *g()* to have the form required by the *atomic* directive.**
2. There is no **guarantee of fairness** in mutual exclusion constructs.
 - For example, the code below may allow one thread to **always** be executing the function *g*, thus preventing other threads from accessing it.

```
while(1) {
    . . .
    # pragma omp critical
    x = g(my_rank);
    . . .
}
```

- This won't happen if the *while* loop **terminates**.
3. It can be dangerous to "**nest**" mutual exclusion constructs.

```
# pragma omp critical
y = f(x);
. . .
double f(double x) {
    # pragma omp critical
    z = g(x); /* z is shared */
    . . .
}
```

- This will cause **deadlock**: a situation in which a group of threads are waiting for each other thread to finish.



- If a thread is executing the first block, it won't be able to **enter the second block**. At the same time, it **will not leave the first block** until it proceeds to the second block.
- One possible solution is use **named critical sections**

```

# pragma omp critical(one)
y = f(x);
. . .
double f(double x) {
#   pragma omp critical(two)
    z = g(x); /* z is global */
    . . .
}
  
```

- But this is not the **ultimate** solution, as deadlocks can occur as following:

Time	Thread <i>u</i>	Thread <i>v</i>
0	Enter crit. sect. one	Enter crit. sect. two
1	Attempt to enter two	Attempt to enter one
2	Block	Block